

Keras -- MLPs on MNIST

In [0]:

```
# if you keras is not using tensorflow as backend set "KERAS_BACKEND=tensorflow" use this command
from keras.utils import np_utils
from keras.datasets import mnist
import seaborn as sns
from keras.initializers import RandomNormal
%matplotlib notebook
```

In [0]:

```
%matplotlib notebook
import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()
```

In [6]:

```
# the data, shuffled and split between train and test sets
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

Downloading data from <https://s3.amazonaws.com/img-datasets/mnist.npz>
11493376/11490434 [=====] - 1s 0us/step

In [7]:

```
print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d, %d)"%(X_train.shape[1], X_train.shape[2]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d, %d)"%(X_test.shape[1], X_test.shape[2]))
```

Number of training examples : 60000 and each image is of shape (28, 28)
Number of training examples : 10000 and each image is of shape (28, 28)

In [0]:

```
# if you observe the input shape its 2 dimensional vector
# for each image we have a (28*28) vector
# we will convert the (28*28) vector into single dimensional vector of 1 * 784

X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
```

In [9]:

```
# after converting the input images from 3d to 2d vectors

print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d)"%(X_train.shape[1]))
print("Number of training examples :", X_test.shape[0], "and each image is of shape (%d)"%(X_test.shape[1]))
```

Number of training examples : 60000 and each image is of shape (784)
Number of training examples : 10000 and each image is of shape (784)

```
# An example data point
print(X_train[0])
```

In [0]:

```
X_train = X_train/255
X_test = X_test/255
```

```
# example data point after normalizing
print(X_train[0])
```

[illegible]

[illegible]

In [0]:

```
# https://keras.io/getting-started/sequential-model-guide/

# The Sequential model is a linear stack of layers.
# you can create a Sequential model by passing a list of layer instances to the constructor:

# model = Sequential([
#     Dense(32, input_shape=(784,)),
#     Activation('relu'),
#     Dense(10),
#     Activation('softmax'),
# ])

# You can also simply add layers via the .add() method:

# model = Sequential()
# model.add(Dense(32, input_dim=784))
# model.add(Activation('relu'))

###

# https://keras.io/layers/core/

# keras.layers.Dense(units, activation=None, use_bias=True, kernel_initializer='glorot_uniform',
# bias_initializer='zeros', kernel_regularizer=None, bias_regularizer=None, activity_regularizer=None,
# kernel_constraint=None, bias_constraint=None)

# Dense implements the operation: output = activation(dot(input, kernel) + bias) where
# activation is the element-wise activation function passed as the activation argument,
# kernel is a weights matrix created by the layer, and
# bias is a bias vector created by the layer (only applicable if use_bias is True).

# output = activation(dot(input, kernel) + bias) => y = activation(WT.X + b)

####

# https://keras.io/activations/

# Activations can either be used through an Activation layer, or through the activation argument supported by all forward layers:

# from keras.layers import Activation, Dense

# model.add(Dense(64))
# model.add(Activation('tanh'))

# This is equivalent to:
# model.add(Dense(64, activation='tanh'))

# there are many activation functions available ex: tanh, relu, softmax

from keras.models import Sequential
from keras.layers import Dense, Activation
```

In [0]:

```
# some model parameters

output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
nb_epoch = 20
```

In [16]:

```

# start building a model
model = Sequential()

# The model needs to know what input shape it should expect.
# For this reason, the first layer in a Sequential model
# (and only the first, because following layers can do automatic shape inference)
# needs to receive information about its input shape.
# you can use input_shape and input_dim to pass the shape of input

# output_dim represent the number of nodes need in that layer
# here we have 10 nodes

model.add(Dense(output_dim, input_dim=input_dim, activation='softmax'))

```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:66: The name tf.get_default_graph is deprecated. Please use tf.compat.v1.get_default_graph instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:541: The name tf.placeholder is deprecated. Please use tf.compat.v1.placeholder instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:4432: The name tf.random_uniform is deprecated. Please use tf.random.uniform instead.

In [17]:

```

# Before training a model, you need to configure the learning process, which is done via
the compile method

# It receives three arguments:
# An optimizer. This could be the string identifier of an existing optimizer , https://ke
ras.io/optimizers/
# A loss function. This is the objective that the model will try to minimize., https://ke
ras.io/losses/
# A list of metrics. For any classification problem you will want to set this to metrics=
['accuracy']. https://keras.io/metrics/

# Note: when using the categorical_crossentropy loss, your targets should be in categoric
al format
# (e.g. if you have 10 classes, the target for each sample should be a 10-dimensional vec
tor that is all-zeros except
# for a 1 at the index corresponding to the class of the sample).

# that is why we converted out labels into vectors

model.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])

# Keras models are trained on Numpy arrays of input data and labels.
# For training a model, you will typically use the fit function

# fit(self, x=None, y=None, batch_size=None, epochs=1, verbose=1, callbacks=None, validat
ion_split=0.0,
# validation_data=None, shuffle=True, class_weight=None, sample_weight=None, initial_epoc
h=0, steps_per_epoch=None,
# validation_steps=None)

# fit() function Trains the model for a fixed number of epochs (iterations on a dataset).

# it returns A History object. Its History.history attribute is a record of training loss
values and
# metrics values at successive epochs, as well as validation loss values and validation m
etrics values (if applicable).

# https://github.com/openai/baselines/issues/20

history1 = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1
, validation_data=(X_test, Y_test))

```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/optimizers.py:793: The name tf.train.Optimizer is deprecated. Please use tf.compat.v1.train.Optimizer instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:3576: The name tf.log is deprecated. Please use tf.math.log instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow_core/python/ops/math_grad.py:1424: where (from tensorflow.python.ops.array_ops) is deprecated and will be removed in a future version.
Instructions for updating:
Use tf.where in 2.0, which has the same broadcast rule as np.where

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:1033: The name tf.assign_add is deprecated. Please use tf.compat.v1.assign_add instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:1020: The name tf.assign is deprecated. Please use tf.compat.v1.assign instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:3005: The name tf.Session is deprecated. Please use tf.compat.v1.Session instead.

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:190: The name tf.get_default_session is deprecated. Please use tf.compat.v1.get_default_session instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:197: The name tf.ConfigProto is deprecated. Please use tf.compat.v1.ConfigProto instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:207: The name tf.global_variables is deprecated. Please use tf.compat.v1.global_variables instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:216: The name tf.is_variable_initialized is deprecated. Please use tf.compat.v1.is_variable_initialized instead.

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:223: The name tf.variables_initializer is deprecated. Please use tf.compat.v1.variables_initializer instead.

60000/60000 [=====] - 11s 188us/step - loss: 1.3074 - acc: 0.6848 - val_loss: 0.8180 - val_acc: 0.8364

Epoch 2/20

60000/60000 [=====] - 1s 24us/step - loss: 0.7217 - acc: 0.8404 - val_loss: 0.6102 - val_acc: 0.8618

Epoch 3/20

60000/60000 [=====] - 1s 25us/step - loss: 0.5905 - acc: 0.8580 - val_loss: 0.5276 - val_acc: 0.8735

Epoch 4/20

60000/60000 [=====] - 2s 25us/step - loss: 0.5279 - acc: 0.8671 - val_loss: 0.4818 - val_acc: 0.8802

Epoch 5/20

60000/60000 [=====] - 2s 27us/step - loss: 0.4898 - acc: 0.8738 - val_loss: 0.4517 - val_acc: 0.8860

Epoch 6/20

60000/60000 [=====] - 2s 25us/step - loss: 0.4637 - acc: 0.8789 - val_loss: 0.4301 - val_acc: 0.8895

Epoch 7/20

60000/60000 [=====] - 2s 27us/step - loss: 0.4443 - acc: 0.8830 - val_loss: 0.4138 - val_acc: 0.8918

Epoch 8/20

60000/60000 [=====] - 1s 23us/step - loss: 0.4292 - acc: 0.8853 - val_loss: 0.4012 - val_acc: 0.8937

Epoch 9/20

60000/60000 [=====] - 2s 25us/step - loss: 0.4170 - acc: 0.8881 - val_loss: 0.3906 - val_acc: 0.8962

Epoch 10/20

```

60000/60000 [=====] - 1s 24us/step - loss: 0.4070 - acc: 0.8899
- val_loss: 0.3817 - val_acc: 0.8979
Epoch 11/20
60000/60000 [=====] - 2s 25us/step - loss: 0.3984 - acc: 0.8918
- val_loss: 0.3745 - val_acc: 0.8985
Epoch 12/20
60000/60000 [=====] - 2s 25us/step - loss: 0.3910 - acc: 0.8933
- val_loss: 0.3678 - val_acc: 0.9004
Epoch 13/20
60000/60000 [=====] - 2s 26us/step - loss: 0.3846 - acc: 0.8950
- val_loss: 0.3624 - val_acc: 0.9016
Epoch 14/20
60000/60000 [=====] - 2s 25us/step - loss: 0.3789 - acc: 0.8961
- val_loss: 0.3573 - val_acc: 0.9032
Epoch 15/20
60000/60000 [=====] - 1s 24us/step - loss: 0.3738 - acc: 0.8973
- val_loss: 0.3532 - val_acc: 0.9045
Epoch 16/20
60000/60000 [=====] - 1s 24us/step - loss: 0.3693 - acc: 0.8980
- val_loss: 0.3492 - val_acc: 0.9055
Epoch 17/20
60000/60000 [=====] - 1s 24us/step - loss: 0.3651 - acc: 0.8995
- val_loss: 0.3454 - val_acc: 0.9071
Epoch 18/20
60000/60000 [=====] - 2s 25us/step - loss: 0.3613 - acc: 0.9001
- val_loss: 0.3421 - val_acc: 0.9075
Epoch 19/20
60000/60000 [=====] - 2s 26us/step - loss: 0.3579 - acc: 0.9011
- val_loss: 0.3389 - val_acc: 0.9093
Epoch 20/20
60000/60000 [=====] - 2s 26us/step - loss: 0.3547 - acc: 0.9019
- val_loss: 0.3363 - val_acc: 0.9093

```

In [0]:

```
%matplotlib inline
```

In [0]:

```

# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.grid()
    fig.canvas.draw()

```

In [20]:

```
print(len(history1.history['val_loss']))
```

20

In [21]:

```

score = model.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])


fig, ax1 = plt.subplots(1, 1)
ax1.set_xlabel('epoch') ; ax1.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x1 = list(range(1, nb_epoch+1))

vy1 = history1.history['val_loss']
ty1 = history1.history['loss']
ax1.plot(x1, vy1, 'b', label="Validation Loss")
ax1.plot(x1, ty1, 'r', label="Train Loss")
plt.grid()
fig.canvas.draw()
plt.show();

```


Test score: 0.3363473850727081
Test accuracy: 0.9093



MLP + Sigmoid activation + SGDOptimizer

In [22]:

```
# Multilayer perceptron

model_sigmoid = Sequential()
model_sigmoid.add(Dense(512, activation='sigmoid', input_shape=(input_dim,)))
model_sigmoid.add(Dense(128, activation='sigmoid'))
model_sigmoid.add(Dense(output_dim, activation='softmax'))

model_sigmoid.summary()
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
=====		
dense_2 (Dense)	(None, 512)	401920
dense_3 (Dense)	(None, 128)	65664
dense_4 (Dense)	(None, 10)	1290
=====		
Total params: 468,874		
Trainable params: 468,874		
Non-trainable params: 0		

In [23]:

```
model_sigmoid.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_sigmoid.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 2s 32us/step - loss: 2.2707 - acc: 0.2212
- val_loss: 2.2289 - val_acc: 0.3618
Epoch 2/20
60000/60000 [=====] - 2s 30us/step - loss: 2.1886 - acc: 0.4307
- val_loss: 2.1386 - val_acc: 0.5356
Epoch 3/20
60000/60000 [=====] - 2s 29us/step - loss: 2.0813 - acc: 0.5460
- val_loss: 2.0068 - val_acc: 0.5419
Epoch 4/20
60000/60000 [=====] - 2s 29us/step - loss: 1.9258 - acc: 0.6055
- val_loss: 1.8199 - val_acc: 0.6352
Epoch 5/20
60000/60000 [=====] - 2s 27us/step - loss: 1.7172 - acc: 0.6523

```

- val_loss: 1.5887 - val_acc: 0.6586
Epoch 6/20
60000/60000 [=====] - 2s 30us/step - loss: 1.4841 - acc: 0.6941
- val_loss: 1.3559 - val_acc: 0.7194
Epoch 7/20
60000/60000 [=====] - 2s 31us/step - loss: 1.2701 - acc: 0.7367
- val_loss: 1.1605 - val_acc: 0.7719
Epoch 8/20
60000/60000 [=====] - 2s 32us/step - loss: 1.0970 - acc: 0.7690
- val_loss: 1.0090 - val_acc: 0.7783
Epoch 9/20
60000/60000 [=====] - 2s 29us/step - loss: 0.9639 - acc: 0.7894
- val_loss: 0.8934 - val_acc: 0.8062
Epoch 10/20
60000/60000 [=====] - 2s 31us/step - loss: 0.8622 - acc: 0.8055
- val_loss: 0.8050 - val_acc: 0.8169
Epoch 11/20
60000/60000 [=====] - 2s 29us/step - loss: 0.7834 - acc: 0.8178
- val_loss: 0.7349 - val_acc: 0.8304
Epoch 12/20
60000/60000 [=====] - 2s 28us/step - loss: 0.7211 - acc: 0.8277
- val_loss: 0.6792 - val_acc: 0.8370
Epoch 13/20
60000/60000 [=====] - 2s 28us/step - loss: 0.6708 - acc: 0.8368
- val_loss: 0.6346 - val_acc: 0.8457
Epoch 14/20
60000/60000 [=====] - 2s 29us/step - loss: 0.6294 - acc: 0.8438
- val_loss: 0.5963 - val_acc: 0.8522
Epoch 15/20
60000/60000 [=====] - 2s 31us/step - loss: 0.5951 - acc: 0.8514
- val_loss: 0.5647 - val_acc: 0.8592
Epoch 16/20
60000/60000 [=====] - 2s 30us/step - loss: 0.5659 - acc: 0.8566
- val_loss: 0.5381 - val_acc: 0.8633
Epoch 17/20
60000/60000 [=====] - 2s 30us/step - loss: 0.5410 - acc: 0.8615
- val_loss: 0.5145 - val_acc: 0.8693
Epoch 18/20
60000/60000 [=====] - 2s 30us/step - loss: 0.5194 - acc: 0.8657
- val_loss: 0.4942 - val_acc: 0.8716
Epoch 19/20
60000/60000 [=====] - 2s 29us/step - loss: 0.5006 - acc: 0.8698
- val_loss: 0.4775 - val_acc: 0.8746
Epoch 20/20
60000/60000 [=====] - 2s 30us/step - loss: 0.4842 - acc: 0.8727
- val_loss: 0.4614 - val_acc: 0.8786

```

In [0]:

```

# list of epoch numbers

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

```

In [24]:

```

x = list(range(1, nb_epoch+1))
score = model_sigmoid.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])


```

```
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch')
ax.set_ylabel('Categorical Crossentropy Loss')

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
fig.canvas.draw()
```

Test score: 0.4614129983663559
 Test accuracy: 0.8786



In [0]:

In [0]:

```
import matplotlib.pyplot as plt
```

In [0]:

```
w_after = model_sigmoid.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)
```

In [27]:

```
w_after = model_sigmoid.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
```

```
plt.xlabel('Output Layer ')
plt.show()
```



MLP + Sigmoid activation + ADAM

In [28]:

```
model_sigmoid = Sequential()
model_sigmoid.add(Dense(512, activation='sigmoid', input_shape=(input_dim,)))
model_sigmoid.add(Dense(128, activation='sigmoid'))
model_sigmoid.add(Dense(output_dim, activation='softmax'))

model_sigmoid.summary()

model_sigmoid.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_sigmoid.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
dense_5 (Dense)	(None, 512)	401920
dense_6 (Dense)	(None, 128)	65664
dense_7 (Dense)	(None, 10)	1290
Total params: 468,874		
Trainable params: 468,874		
Non-trainable params: 0		

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 2s 40us/step - loss: 0.5336 - acc: 0.8603
- val_loss: 0.2545 - val_acc: 0.9255
Epoch 2/20
60000/60000 [=====] - 2s 35us/step - loss: 0.2227 - acc: 0.9348
- val_loss: 0.1845 - val_acc: 0.9438
Epoch 3/20
60000/60000 [=====] - 2s 36us/step - loss: 0.1642 - acc: 0.9505
- val_loss: 0.1497 - val_acc: 0.9532
Epoch 4/20
60000/60000 [=====] - 2s 35us/step - loss: 0.1257 - acc: 0.9629
- val_loss: 0.1179 - val_acc: 0.9645
Epoch 5/20
60000/60000 [=====] - 2s 35us/step - loss: 0.0986 - acc: 0.9709
- val_loss: 0.1013 - val_acc: 0.9690
Epoch 6/20
60000/60000 [=====] - 2s 36us/step - loss: 0.0803 - acc: 0.9763
- val_loss: 0.0919 - val_acc: 0.9718
Epoch 7/20
```

```

Epoch 7/20
60000/60000 [=====] - 2s 36us/step - loss: 0.0641 - acc: 0.9813
- val_loss: 0.0826 - val_acc: 0.9736
Epoch 8/20
60000/60000 [=====] - 2s 37us/step - loss: 0.0519 - acc: 0.9848
- val_loss: 0.0752 - val_acc: 0.9771
Epoch 9/20
60000/60000 [=====] - 2s 38us/step - loss: 0.0435 - acc: 0.9873
- val_loss: 0.0713 - val_acc: 0.9784
Epoch 10/20
60000/60000 [=====] - 2s 33us/step - loss: 0.0352 - acc: 0.9899
- val_loss: 0.0759 - val_acc: 0.9762
Epoch 11/20
60000/60000 [=====] - 2s 35us/step - loss: 0.0283 - acc: 0.9920
- val_loss: 0.0672 - val_acc: 0.9802
Epoch 12/20
60000/60000 [=====] - 2s 36us/step - loss: 0.0217 - acc: 0.9947
- val_loss: 0.0616 - val_acc: 0.9820
Epoch 13/20
60000/60000 [=====] - 2s 37us/step - loss: 0.0181 - acc: 0.9956
- val_loss: 0.0662 - val_acc: 0.9799
Epoch 14/20
60000/60000 [=====] - 2s 35us/step - loss: 0.0142 - acc: 0.9968
- val_loss: 0.0753 - val_acc: 0.9780
Epoch 15/20
60000/60000 [=====] - 2s 35us/step - loss: 0.0123 - acc: 0.9968
- val_loss: 0.0656 - val_acc: 0.9805
Epoch 16/20
60000/60000 [=====] - 2s 34us/step - loss: 0.0090 - acc: 0.9980
- val_loss: 0.0703 - val_acc: 0.9806
Epoch 17/20
60000/60000 [=====] - 2s 35us/step - loss: 0.0071 - acc: 0.9986
- val_loss: 0.0742 - val_acc: 0.9789
Epoch 18/20
60000/60000 [=====] - 2s 35us/step - loss: 0.0071 - acc: 0.9982
- val_loss: 0.0716 - val_acc: 0.9804
Epoch 19/20
60000/60000 [=====] - 2s 34us/step - loss: 0.0053 - acc: 0.9990
- val_loss: 0.0666 - val_acc: 0.9825
Epoch 20/20
60000/60000 [=====] - 2s 34us/step - loss: 0.0035 - acc: 0.9993
- val_loss: 0.0725 - val_acc: 0.9815

```

In [29]:

```

score = model_sigmoid.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.07249096132973064

Test accuracy: 0.9815



In [30]:


```
w_after = model_sigmoid.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + ReLU +SGD

In [31]:

```
# Multilayer perceptron
```

```
# https://arxiv.org/pdf/1707.09725.pdf#page=95
# for relu layers
# If we sample weights from a normal distribution  $N(0,\sigma)$  we satisfy this condition with  $\sigma = \sqrt{2/(ni)}$ .
# h1 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.062 \Rightarrow N(0,\sigma) = N(0,0.062)$ 
# h2 =>  $\sigma = \sqrt{2/(fan\_in)} = 0.125 \Rightarrow N(0,\sigma) = N(0,0.125)$ 
# out =>  $\sigma = \sqrt{2/(fan\_in+1)} = 0.120 \Rightarrow N(0,\sigma) = N(0,0.120)$ 
```

```
model_relu = Sequential()
model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
model_relu.add(Dense(output_dim, activation='softmax'))

model_relu.summary()
```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:4409: The name tf.random_normal is deprecated. Please use tf.random.normal instead.

Model: "sequential_4"

Layer (type)	Output Shape	Param #
dense_8 (Dense)	(None, 512)	401920
dense_9 (Dense)	(None, 128)	65664
dense_10 (Dense)	(None, 10)	1290
Total params: 468,874		
Trainable params: 468,874		
Non-trainable params: 0		

In [32]:

```
model_relu.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

```
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 2s 31us/step - loss: 0.7522 - acc: 0.7912
- val_loss: 0.3851 - val_acc: 0.8953
Epoch 2/20
60000/60000 [=====] - 2s 32us/step - loss: 0.3552 - acc: 0.8993
- val_loss: 0.2979 - val_acc: 0.9158
Epoch 3/20
60000/60000 [=====] - 2s 31us/step - loss: 0.2924 - acc: 0.9173
- val_loss: 0.2600 - val_acc: 0.9264
Epoch 4/20
60000/60000 [=====] - 2s 32us/step - loss: 0.2585 - acc: 0.9269
- val_loss: 0.2405 - val_acc: 0.9308
Epoch 5/20
60000/60000 [=====] - 2s 29us/step - loss: 0.2352 - acc: 0.9335
- val_loss: 0.2202 - val_acc: 0.9365
Epoch 6/20
60000/60000 [=====] - 2s 31us/step - loss: 0.2171 - acc: 0.9380
- val_loss: 0.2085 - val_acc: 0.9404
Epoch 7/20
60000/60000 [=====] - 2s 33us/step - loss: 0.2023 - acc: 0.9427
- val_loss: 0.1957 - val_acc: 0.9426
Epoch 8/20
60000/60000 [=====] - 2s 32us/step - loss: 0.1903 - acc: 0.9455
- val_loss: 0.1857 - val_acc: 0.9452
Epoch 9/20
60000/60000 [=====] - 2s 32us/step - loss: 0.1795 - acc: 0.9487
- val_loss: 0.1765 - val_acc: 0.9487
```

```

Epoch 10/20
60000/60000 [=====] - 2s 31us/step - loss: 0.1699 - acc: 0.9512
- val_loss: 0.1699 - val_acc: 0.9502
Epoch 11/20
60000/60000 [=====] - 2s 31us/step - loss: 0.1616 - acc: 0.9540
- val_loss: 0.1640 - val_acc: 0.9506
Epoch 12/20
60000/60000 [=====] - 2s 31us/step - loss: 0.1539 - acc: 0.9566
- val_loss: 0.1583 - val_acc: 0.9531
Epoch 13/20
60000/60000 [=====] - 2s 28us/step - loss: 0.1471 - acc: 0.9583
- val_loss: 0.1519 - val_acc: 0.9546
Epoch 14/20
60000/60000 [=====] - 2s 28us/step - loss: 0.1407 - acc: 0.9599
- val_loss: 0.1470 - val_acc: 0.9571
Epoch 15/20
60000/60000 [=====] - 2s 28us/step - loss: 0.1347 - acc: 0.9622
- val_loss: 0.1421 - val_acc: 0.9572
Epoch 16/20
60000/60000 [=====] - 2s 28us/step - loss: 0.1294 - acc: 0.9639
- val_loss: 0.1390 - val_acc: 0.9580
Epoch 17/20
60000/60000 [=====] - 2s 28us/step - loss: 0.1245 - acc: 0.9650
- val_loss: 0.1366 - val_acc: 0.9603
Epoch 18/20
60000/60000 [=====] - 2s 33us/step - loss: 0.1198 - acc: 0.9663
- val_loss: 0.1305 - val_acc: 0.9620
Epoch 19/20
60000/60000 [=====] - 2s 33us/step - loss: 0.1153 - acc: 0.9681
- val_loss: 0.1283 - val_acc: 0.9613
Epoch 20/20
60000/60000 [=====] - 2s 30us/step - loss: 0.1112 - acc: 0.9693
- val_loss: 0.1261 - val_acc: 0.9617

```

In [33]:

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verb
ose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epoc
hs


vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.12610071545392273

Test accuracy: 0.9617





In [34]:


```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + ReLU + ADAM

In [35]:

```
model_relu = Sequential()
model_relu.add(Dense(512, activation='relu', input_shape=(input_dim,), kernel_initializer
=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
model_relu.add(Dense(128, activation='relu', kernel_initializer=RandomNormal(mean=0.0, st
ddev=0.125, seed=None)) )
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())

model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy
'])
```

```
history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Model: "sequential_5"

Layer (type)	Output Shape	Param #
dense_11 (Dense)	(None, 512)	401920
dense_12 (Dense)	(None, 128)	65664
dense_13 (Dense)	(None, 10)	1290

Total params: 468,874
Trainable params: 468,874
Non-trainable params: 0

```
None
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 2s 41us/step - loss: 0.2252 - acc: 0.9319
- val_loss: 0.1124 - val_acc: 0.9659
Epoch 2/20
60000/60000 [=====] - 2s 36us/step - loss: 0.0852 - acc: 0.9742
- val_loss: 0.0806 - val_acc: 0.9741
Epoch 3/20
60000/60000 [=====] - 2s 34us/step - loss: 0.0533 - acc: 0.9832
- val_loss: 0.0892 - val_acc: 0.9715
Epoch 4/20
60000/60000 [=====] - 2s 34us/step - loss: 0.0365 - acc: 0.9883
- val_loss: 0.0703 - val_acc: 0.9764
Epoch 5/20
60000/60000 [=====] - 2s 38us/step - loss: 0.0250 - acc: 0.9919
- val_loss: 0.0650 - val_acc: 0.9814
Epoch 6/20
60000/60000 [=====] - 2s 36us/step - loss: 0.0201 - acc: 0.9935
- val_loss: 0.0704 - val_acc: 0.9804
Epoch 7/20
60000/60000 [=====] - 2s 37us/step - loss: 0.0169 - acc: 0.9942
- val_loss: 0.0720 - val_acc: 0.9790
Epoch 8/20
60000/60000 [=====] - 2s 36us/step - loss: 0.0151 - acc: 0.9952
- val_loss: 0.0727 - val_acc: 0.9800
Epoch 9/20
60000/60000 [=====] - 2s 36us/step - loss: 0.0129 - acc: 0.9957
- val_loss: 0.0751 - val_acc: 0.9805
Epoch 10/20
60000/60000 [=====] - 2s 33us/step - loss: 0.0131 - acc: 0.9956
- val_loss: 0.1060 - val_acc: 0.9747
Epoch 11/20
60000/60000 [=====] - 2s 38us/step - loss: 0.0131 - acc: 0.9954
- val_loss: 0.0898 - val_acc: 0.9775
Epoch 12/20
60000/60000 [=====] - 2s 34us/step - loss: 0.0099 - acc: 0.9969
- val_loss: 0.0870 - val_acc: 0.9798
Epoch 13/20
60000/60000 [=====] - 2s 37us/step - loss: 0.0085 - acc: 0.9970
- val_loss: 0.0841 - val_acc: 0.9793
Epoch 14/20
60000/60000 [=====] - 2s 35us/step - loss: 0.0110 - acc: 0.9961
- val_loss: 0.0866 - val_acc: 0.9813
Epoch 15/20
60000/60000 [=====] - 2s 37us/step - loss: 0.0053 - acc: 0.9984
- val_loss: 0.0862 - val_acc: 0.9809
Epoch 16/20
60000/60000 [=====] - 2s 36us/step - loss: 0.0081 - acc: 0.9972
- val_loss: 0.1060 - val_acc: 0.9780
Epoch 17/20
60000/60000 [=====] - 2s 36us/step - loss: 0.0104 - acc: 0.9966
- val_loss: 0.0843 - val_acc: 0.9818
Epoch 18/20
```

```
60000/60000 [=====] - 2s 37us/step - loss: 0.0104 - acc: 0.9964
- val_loss: 0.1003 - val_acc: 0.9798
Epoch 19/20
60000/60000 [=====] - 2s 33us/step - loss: 0.0039 - acc: 0.9987
- val_loss: 0.0854 - val_acc: 0.9833
Epoch 20/20
60000/60000 [=====] - 2s 35us/step - loss: 0.0057 - acc: 0.9983
- val_loss: 0.0834 - val_acc: 0.9843
```

In [36]:

```
score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verb
ose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epoc
hs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.08342993713870042

Test accuracy: 0.9843



In [37]:

```
w_after = model_relu.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
```

```
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + Batch-Norm on hidden Layers + AdamOptimizer

In [38]:

```
# Multilayer perceptron

# https://intoli.com/blog/neural-network-initialization/
# If we sample weights from a normal distribution  $N(0, \sigma)$  we satisfy this condition with  $\sigma = \sqrt{2/(n_i + n_{i+1})}$ .
# h1 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.039 \Rightarrow N(0, \sigma) = N(0, 0.039)$ 
# h2 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.055 \Rightarrow N(0, \sigma) = N(0, 0.055)$ 
# h1 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.120 \Rightarrow N(0, \sigma) = N(0, 0.120)$ 

from keras.layers.normalization import BatchNormalization

model_batch = Sequential()

model_batch.add(Dense(512, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(128, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))

model_batch.summary()
```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:148: The name tf.placeholder_with_default is deprecated. Please use tf.compat.v1.placeholder_with_default instead.

Model: "sequential_6"

Layer (type)	Output shape	Param #
dense_14 (Dense)	(None, 512)	401920
batch_normalization_1 (Batch Normalization)	(None, 512)	2048
dense_15 (Dense)	(None, 128)	65664
batch_normalization_2 (Batch Normalization)	(None, 128)	512
dense_16 (Dense)	(None, 10)	1290
Total params: 471,434		
Trainable params: 470,154		
Non-trainable params: 1,280		

In [39]:

```
model_batch.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model_batch.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 4s 68us/step - loss: 0.2960 - acc: 0.9113
- val_loss: 0.2120 - val_acc: 0.9373
Epoch 2/20
60000/60000 [=====] - 4s 59us/step - loss: 0.1723 - acc: 0.9495
- val_loss: 0.1736 - val_acc: 0.9471
Epoch 3/20
60000/60000 [=====] - 4s 62us/step - loss: 0.1349 - acc: 0.9607
- val_loss: 0.1430 - val_acc: 0.9574
Epoch 4/20
60000/60000 [=====] - 3s 55us/step - loss: 0.1106 - acc: 0.9673
- val_loss: 0.1429 - val_acc: 0.9568
Epoch 5/20
60000/60000 [=====] - 3s 58us/step - loss: 0.0953 - acc: 0.9712
- val_loss: 0.1245 - val_acc: 0.9624
Epoch 6/20
60000/60000 [=====] - 3s 56us/step - loss: 0.0797 - acc: 0.9758
- val_loss: 0.1142 - val_acc: 0.9648
Epoch 7/20
60000/60000 [=====] - 4s 61us/step - loss: 0.0694 - acc: 0.9791
- val_loss: 0.1101 - val_acc: 0.9648
Epoch 8/20
60000/60000 [=====] - 4s 60us/step - loss: 0.0581 - acc: 0.9823
- val_loss: 0.1061 - val_acc: 0.9697
Epoch 9/20
60000/60000 [=====] - 4s 59us/step - loss: 0.0519 - acc: 0.9839
- val_loss: 0.1064 - val_acc: 0.9675
Epoch 10/20
60000/60000 [=====] - 3s 55us/step - loss: 0.0442 - acc: 0.9860
- val_loss: 0.1013 - val_acc: 0.9693
Epoch 11/20
60000/60000 [=====] - 4s 60us/step - loss: 0.0390 - acc: 0.9878
- val_loss: 0.1058 - val_acc: 0.9690
Epoch 12/20
60000/60000 [=====] - 4s 61us/step - loss: 0.0346 - acc: 0.9891
- val_loss: 0.0983 - val_acc: 0.9709
Epoch 13/20
60000/60000 [=====] - 3s 58us/step - loss: 0.0305 - acc: 0.9898
- val_loss: 0.1001 - val_acc: 0.9730
Epoch 14/20
60000/60000 [=====] - 4s 61us/step - loss: 0.0282 - acc: 0.9910
- val_loss: 0.0978 - val_acc: 0.9735
Epoch 15/20
60000/60000 [=====] - 3s 57us/step - loss: 0.0241 - acc: 0.9928
- val_loss: 0.1007 - val_acc: 0.9724
Epoch 16/20
60000/60000 [=====] - 3s 57us/step - loss: 0.0188 - acc: 0.9938
- val_loss: 0.1000 - val_acc: 0.9724
```

```

60000/60000 [=====] - 3s 57us/step - loss: 0.0199 - acc: 0.9938
- val_loss: 0.0976 - val_acc: 0.9736
Epoch 17/20
60000/60000 [=====] - 3s 57us/step - loss: 0.0187 - acc: 0.9937
- val_loss: 0.1039 - val_acc: 0.9731
Epoch 18/20
60000/60000 [=====] - 4s 60us/step - loss: 0.0201 - acc: 0.9930
- val_loss: 0.1009 - val_acc: 0.9730
Epoch 19/20
60000/60000 [=====] - 4s 60us/step - loss: 0.0163 - acc: 0.9951
- val_loss: 0.0933 - val_acc: 0.9766
Epoch 20/20
60000/60000 [=====] - 4s 64us/step - loss: 0.0156 - acc: 0.9950
- val_loss: 0.1062 - val_acc: 0.9737

```

In [40]:

```

score = model_batch.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.10622683503271255

Test accuracy: 0.9737



In [41]:

```

w_after = model_batch.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

```

```
fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



5. MLP + Dropout + AdamOptimizer

In [42]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras

from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(512, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(128, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:3733: calling dropout (from tensorflow.python.ops.nn_ops) with keep_prob is deprecated and will be removed in a future version.

Instructions for updating:

Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 - keep_prob`.

Model: "sequential_7"

Layer (type)	Output Shape	Param #
dense_17 (Dense)	(None, 512)	401920
batch_normalization_3 (Batch Normalization)	(None, 512)	2048
dropout_1 (Dropout)	(None, 512)	0
dense_18 (Dense)	(None, 128)	65664
batch_normalization_4 (Batch Normalization)	(None, 128)	512
dropout_2 (Dropout)	(None, 128)	0
dense_19 (Dense)	(None, 10)	1290
Total params: 471,434		
Trainable params: 470,154		
Non-trainable params: 1,280		

In [43]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 5s 77us/step - loss: 0.6715 - acc: 0.7927
- val_loss: 0.2873 - val_acc: 0.9125
Epoch 2/20
60000/60000 [=====] - 4s 62us/step - loss: 0.4303 - acc: 0.8689
- val_loss: 0.2536 - val_acc: 0.9257
Epoch 3/20
60000/60000 [=====] - 4s 64us/step - loss: 0.3826 - acc: 0.8837
- val_loss: 0.2402 - val_acc: 0.9292
Epoch 4/20
60000/60000 [=====] - 4s 60us/step - loss: 0.3532 - acc: 0.8934
- val_loss: 0.2198 - val_acc: 0.9346
Epoch 5/20
60000/60000 [=====] - 4s 59us/step - loss: 0.3368 - acc: 0.8979
- val_loss: 0.2066 - val_acc: 0.9399
Epoch 6/20
60000/60000 [=====] - 4s 61us/step - loss: 0.3244 - acc: 0.9024
- val_loss: 0.1965 - val_acc: 0.9419
Epoch 7/20
60000/60000 [=====] - 4s 66us/step - loss: 0.3088 - acc: 0.9073
- val_loss: 0.1898 - val_acc: 0.9443
Epoch 8/20
60000/60000 [=====] - 3s 57us/step - loss: 0.2899 - acc: 0.9126
- val_loss: 0.1807 - val_acc: 0.9448
Epoch 9/20
60000/60000 [=====] - 4s 66us/step - loss: 0.2821 - acc: 0.9149
- val_loss: 0.1770 - val_acc: 0.9481
Epoch 10/20
60000/60000 [=====] - 4s 61us/step - loss: 0.2687 - acc: 0.9208
- val_loss: 0.1653 - val_acc: 0.9505
Epoch 11/20
60000/60000 [=====] - 4s 63us/step - loss: 0.2623 - acc: 0.9203
- val_loss: 0.1628 - val_acc: 0.9495
Epoch 12/20
60000/60000 [=====] - 3s 58us/step - loss: 0.2544 - acc: 0.9241
- val_loss: 0.1512 - val_acc: 0.9556
Epoch 13/20
60000/60000 [=====] - 4s 65us/step - loss: 0.2346 - acc: 0.9299
- val_loss: 0.1439 - val_acc: 0.9576
Epoch 14/20
60000/60000 [=====] - 4s 65us/step - loss: 0.2281 - acc: 0.9305
- val_loss: 0.1379 - val_acc: 0.9587
Epoch 15/20
```



```
Epoch 16/20
60000/60000 [=====] - 4s 62us/step - loss: 0.2179 - acc: 0.9335
- val_loss: 0.1325 - val_acc: 0.9601
Epoch 16/20
60000/60000 [=====] - 4s 60us/step - loss: 0.2102 - acc: 0.9363
- val_loss: 0.1243 - val_acc: 0.9633
Epoch 17/20
60000/60000 [=====] - 4s 61us/step - loss: 0.1977 - acc: 0.9408
- val_loss: 0.1192 - val_acc: 0.9651
Epoch 18/20
60000/60000 [=====] - 4s 59us/step - loss: 0.1924 - acc: 0.9423
- val_loss: 0.1176 - val_acc: 0.9645
Epoch 19/20
60000/60000 [=====] - 4s 62us/step - loss: 0.1840 - acc: 0.9444
- val_loss: 0.1107 - val_acc: 0.9661
Epoch 20/20
60000/60000 [=====] - 4s 60us/step - loss: 0.1724 - acc: 0.9487
- val_loss: 0.1100 - val_acc: 0.9664
```

In [44]:

```
score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verb
ose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epoc
hs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.10995881164651364

Test accuracy: 0.9664



In [45]:


```
w_after = model_drop.get_weights()
```

```
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)
```

```
fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



Hyper-parameter tuning of Keras models using Sklearn

In [0]:

```
from keras.optimizers import Adam, RMSprop, SGD
def best_hyperparameters(activ):

    model = Sequential()
    model.add(Dense(512, activation=activ, input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
    model.add(Dense(128, activation=activ, kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)) )
    model.add(Dense(output_dim, activation='softmax'))

    model.compile(loss='categorical_crossentropy', metrics=['accuracy'], optimizer='adam')

    return model
```

In [0]:

```
# https://machinelearningmastery.com/grid-search-hyperparameters-deep-learning-models-python-keras/

activ = ['sigmoid', 'relu']

from keras.wrappers.scikit_learn import KerasClassifier
from sklearn.model_selection import GridSearchCV
```

```

model = KerasClassifier(build_fn=best_hyperparameters, epochs=nb_epoch, batch_size=batch_size, verbose=0)
param_grid = dict(activ=activ)

# if you are using CPU
# grid = GridSearchCV(estimator=model, param_grid=param_grid, n_jobs=-1)
# if you are using GPU dont use the n_jobs parameter

grid = GridSearchCV(estimator=model, param_grid=param_grid)
grid_result = grid.fit(X_train, Y_train)

```

In [48]:

```

print("Best: %f using %s" % (grid_result.best_score_, grid_result.best_params_))
means = grid_result.cv_results_['mean_test_score']
stds = grid_result.cv_results_['std_test_score']
params = grid_result.cv_results_['params']
for mean, stdev, param in zip(means, stds, params):
    print("%f (%f) with: %r" % (mean, stdev, param))

```

```

Best: 0.978583 using {'activ': 'relu'}
0.977783 (0.001954) with: {'activ': 'sigmoid'}
0.978583 (0.001541) with: {'activ': 'relu'}

```

In [0]:

In [0]:

In [0]:

Two Hidden layer Architecture

In [0]:

Using RELU Activation and Adam Optimizer

In [49]:

```

# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras

from keras.layers import Dropout

model1 = Sequential()

model1.add(Dense(352, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))

model1.add(Dense(52, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )

model1.add(Dense(output_dim, activation='softmax'))

model1.summary()

```

Model: "sequential_19"

Layer (type)	Output Shape	Param #
dense_53 (Dense)	(None, 352)	276320
dense_54 (Dense)	(None, 52)	18356
dense_55 (Dense)	(None, 10)	530
Total params: 295,206		
Trainable params: 295,206		
Non-trainable params: 0		

In [50]:

```
model1.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model1.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1,
                    validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 3s 58us/step - loss: 0.2520 - acc: 0.9256
- val_loss: 0.1338 - val_acc: 0.9580
Epoch 2/20
60000/60000 [=====] - 2s 35us/step - loss: 0.0972 - acc: 0.9699
- val_loss: 0.1128 - val_acc: 0.9665
Epoch 3/20
60000/60000 [=====] - 2s 42us/step - loss: 0.0652 - acc: 0.9801
- val_loss: 0.0927 - val_acc: 0.9717
Epoch 4/20
60000/60000 [=====] - 2s 36us/step - loss: 0.0484 - acc: 0.9845
- val_loss: 0.0953 - val_acc: 0.9694
Epoch 5/20
60000/60000 [=====] - 2s 35us/step - loss: 0.0355 - acc: 0.9885
- val_loss: 0.0854 - val_acc: 0.9766
Epoch 6/20
60000/60000 [=====] - 2s 38us/step - loss: 0.0306 - acc: 0.9901
- val_loss: 0.1109 - val_acc: 0.9700
Epoch 7/20
60000/60000 [=====] - 2s 37us/step - loss: 0.0249 - acc: 0.9919
- val_loss: 0.0891 - val_acc: 0.9743
Epoch 8/20
60000/60000 [=====] - 2s 33us/step - loss: 0.0199 - acc: 0.9933
- val_loss: 0.0971 - val_acc: 0.9757
Epoch 9/20
60000/60000 [=====] - 2s 33us/step - loss: 0.0207 - acc: 0.9930
- val_loss: 0.0901 - val_acc: 0.9752
Epoch 10/20
60000/60000 [=====] - 2s 35us/step - loss: 0.0156 - acc: 0.9950
- val_loss: 0.1010 - val_acc: 0.9753
Epoch 11/20
60000/60000 [=====] - 2s 35us/step - loss: 0.0153 - acc: 0.9948
- val_loss: 0.1056 - val_acc: 0.9764
Epoch 12/20
60000/60000 [=====] - 2s 40us/step - loss: 0.0137 - acc: 0.9953
- val_loss: 0.0913 - val_acc: 0.9780
Epoch 13/20
60000/60000 [=====] - 2s 42us/step - loss: 0.0137 - acc: 0.9954
- val_loss: 0.0965 - val_acc: 0.9779
Epoch 14/20
60000/60000 [=====] - 2s 37us/step - loss: 0.0132 - acc: 0.9956
- val_loss: 0.1077 - val_acc: 0.9769
Epoch 15/20
60000/60000 [=====] - 2s 36us/step - loss: 0.0184 - acc: 0.9938
- val_loss: 0.0926 - val_acc: 0.9795
Epoch 16/20
60000/60000 [=====] - 2s 39us/step - loss: 0.0106 - acc: 0.9966
- val_loss: 0.1087 - val_acc: 0.9790
Epoch 17/20
60000/60000 [=====] - 2s 38us/step - loss: 0.0085 - acc: 0.9974
```

```
- val_loss: 0.1010 - val_acc: 0.9792
Epoch 18/20
60000/60000 [=====] - 2s 36us/step - loss: 0.0141 - acc: 0.9957
- val_loss: 0.1195 - val_acc: 0.9775
Epoch 19/20
60000/60000 [=====] - 2s 36us/step - loss: 0.0080 - acc: 0.9974
- val_loss: 0.0957 - val_acc: 0.9801
Epoch 20/20
60000/60000 [=====] - 2s 38us/step - loss: 0.0041 - acc: 0.9987
- val_loss: 0.1067 - val_acc: 0.9769
```

In [51]:

```
score = model1.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.10668754959471503

Test accuracy: 0.9769



In [52]:

```
w_after = model1.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
```

```
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w, color='b')
plt.xlabel('Hidden Layer 1')
```

```
plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2')
```

```
plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w, color='y')
plt.xlabel('Output Layer')
plt.show()
```



In [0]:

In [0]:

In [0]:

With Dropout

In [53]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras

from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(352, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_drop.add(Dropout(0.5))

model_drop.add(Dense(52, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)))
model_drop.add(Dropout(0.7))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

WARNING:tensorflow:Large dropout rate: 0.7 (>0.5). In TensorFlow 2.x, dropout() uses drop

out rate instead of keep_prob. Please ensure that this is intended.
Model: "sequential_20"

Layer (type)	Output Shape	Param #
=====		
dense_56 (Dense)	(None, 352)	276320

dropout_3 (Dropout)	(None, 352)	0

dense_57 (Dense)	(None, 52)	18356

dropout_4 (Dropout)	(None, 52)	0

dense_58 (Dense)	(None, 10)	530
=====		
Total params: 295,206		
Trainable params: 295,206		
Non-trainable params: 0		

In [54]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 4s 60us/step - loss: 1.3294 - acc: 0.5665
- val_loss: 0.3412 - val_acc: 0.9163

Epoch 2/20

60000/60000 [=====] - 2s 37us/step - loss: 0.7147 - acc: 0.7711
- val_loss: 0.2594 - val_acc: 0.9316

Epoch 3/20

60000/60000 [=====] - 2s 40us/step - loss: 0.5678 - acc: 0.8272
- val_loss: 0.2178 - val_acc: 0.9427

Epoch 4/20

60000/60000 [=====] - 2s 40us/step - loss: 0.4968 - acc: 0.8510
- val_loss: 0.2039 - val_acc: 0.9473

Epoch 5/20

60000/60000 [=====] - 2s 40us/step - loss: 0.4382 - acc: 0.8706
- val_loss: 0.1824 - val_acc: 0.9531

Epoch 6/20

60000/60000 [=====] - 2s 42us/step - loss: 0.3978 - acc: 0.8834
- val_loss: 0.1722 - val_acc: 0.9550

Epoch 7/20

60000/60000 [=====] - 2s 36us/step - loss: 0.3694 - acc: 0.8918
- val_loss: 0.1599 - val_acc: 0.9575

Epoch 8/20

60000/60000 [=====] - 3s 42us/step - loss: 0.3429 - acc: 0.8998
- val_loss: 0.1586 - val_acc: 0.9588

Epoch 9/20

60000/60000 [=====] - 2s 40us/step - loss: 0.3345 - acc: 0.9029
- val_loss: 0.1439 - val_acc: 0.9614

Epoch 10/20

60000/60000 [=====] - 2s 38us/step - loss: 0.3093 - acc: 0.9097
- val_loss: 0.1484 - val_acc: 0.9633

Epoch 11/20

60000/60000 [=====] - 2s 38us/step - loss: 0.2965 - acc: 0.9162
- val_loss: 0.1440 - val_acc: 0.9642

Epoch 12/20

60000/60000 [=====] - 2s 37us/step - loss: 0.2857 - acc: 0.9177
- val_loss: 0.1309 - val_acc: 0.9665

Epoch 13/20

60000/60000 [=====] - 2s 40us/step - loss: 0.2770 - acc: 0.9205
- val_loss: 0.1416 - val_acc: 0.9662

Epoch 14/20

60000/60000 [=====] - 2s 40us/step - loss: 0.2674 - acc: 0.9227
- val_loss: 0.1549 - val_acc: 0.9646

```

val_loss: 0.1351 - val_acc: 0.9676
Epoch 15/20
60000/60000 [=====] - 2s 36us/step - loss: 0.2615 - acc: 0.9252
- val_loss: 0.1351 - val_acc: 0.9676
Epoch 16/20
60000/60000 [=====] - 2s 39us/step - loss: 0.2521 - acc: 0.9278
- val_loss: 0.1285 - val_acc: 0.9686
Epoch 17/20
60000/60000 [=====] - 3s 44us/step - loss: 0.2452 - acc: 0.9278
- val_loss: 0.1281 - val_acc: 0.9691
Epoch 18/20
60000/60000 [=====] - 2s 39us/step - loss: 0.2397 - acc: 0.9307
- val_loss: 0.1285 - val_acc: 0.9693
Epoch 19/20
60000/60000 [=====] - 2s 39us/step - loss: 0.2326 - acc: 0.9328
- val_loss: 0.1348 - val_acc: 0.9710
Epoch 20/20
60000/60000 [=====] - 2s 39us/step - loss: 0.2283 - acc: 0.9345
- val_loss: 0.1325 - val_acc: 0.9708

```

In [55]:

```

score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verb
ose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epoc
hs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.13253417925792746

Test accuracy: 0.9708



In [56]:

```
w_after = model_drop.get_weights()
```

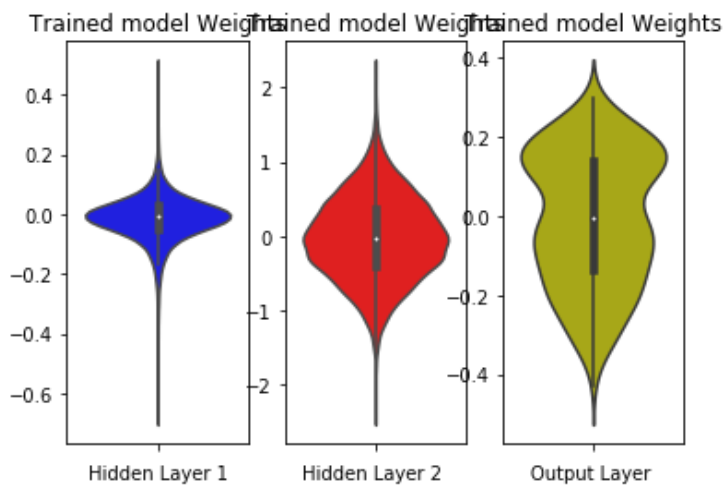


```
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)
```

```
fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



In [0]:

With Batch Normalization

In [0]:

In [57]:

```
# Multilayer perceptron

# https://intoli.com/blog/neural-network-initialization/
# If we sample weights from a normal distribution  $N(0,\sigma)$  we satisfy this condition with  $\sigma = \sqrt{2/(n_i+n_{i+1})}$ .
# h1 =>  $\sigma = \sqrt{2/(n_i+n_{i+1})} = 0.039 \Rightarrow N(0,\sigma) = N(0,0.039)$ 
# h2 =>  $\sigma = \sqrt{2/(n_i+n_{i+1})} = 0.055 \Rightarrow N(0,\sigma) = N(0,0.055)$ 
# h1 =>  $\sigma = \sqrt{2/(n_i+n_{i+1})} = 0.120 \Rightarrow N(0,\sigma) = N(0,0.120)$ 

from keras.layers.normalization import BatchNormalization

model_batch = Sequential()

model_batch.add(Dense(352, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_batch.add(BatchNormalization())
```

```
model_batch.add(Dense(52, activation='relu', kernel_initializer=RandomNormal(mean=0.0, s
tddev=0.55, seed=None)) )
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))

model_batch.summary()
```

Model: "sequential_21"

Layer (type)	Output Shape	Param #
dense_59 (Dense)	(None, 352)	276320
batch_normalization_5 (Batch Normalization)	(None, 352)	1408
dense_60 (Dense)	(None, 52)	18356
batch_normalization_6 (Batch Normalization)	(None, 52)	208
dense_61 (Dense)	(None, 10)	530
Total params: 296,822		
Trainable params: 296,014		
Non-trainable params: 808		

In [58]:

```
model_batch.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_batch.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20
60000/60000 [=====] - 5s 87us/step - loss: 0.2292 - acc: 0.9376
- val_loss: 0.1217 - val_acc: 0.9638

Epoch 2/20
60000/60000 [=====] - 4s 63us/step - loss: 0.0858 - acc: 0.9756
- val_loss: 0.0932 - val_acc: 0.9694

Epoch 3/20
60000/60000 [=====] - 3s 58us/step - loss: 0.0542 - acc: 0.9838
- val_loss: 0.0933 - val_acc: 0.9707

Epoch 4/20
60000/60000 [=====] - 3s 58us/step - loss: 0.0384 - acc: 0.9882
- val_loss: 0.0948 - val_acc: 0.9696

Epoch 5/20
60000/60000 [=====] - 4s 59us/step - loss: 0.0294 - acc: 0.9907
- val_loss: 0.0742 - val_acc: 0.9777

Epoch 6/20
60000/60000 [=====] - 4s 60us/step - loss: 0.0233 - acc: 0.9927
- val_loss: 0.0752 - val_acc: 0.9781

Epoch 7/20
60000/60000 [=====] - 3s 57us/step - loss: 0.0214 - acc: 0.9930
- val_loss: 0.0866 - val_acc: 0.9742

Epoch 8/20
60000/60000 [=====] - 4s 62us/step - loss: 0.0183 - acc: 0.9942
- val_loss: 0.0771 - val_acc: 0.9781

Epoch 9/20
60000/60000 [=====] - 4s 60us/step - loss: 0.0125 - acc: 0.9964
- val_loss: 0.0759 - val_acc: 0.9776

Epoch 10/20
60000/60000 [=====] - 3s 52us/step - loss: 0.0116 - acc: 0.9966
- val_loss: 0.0798 - val_acc: 0.9784

Epoch 11/20
60000/60000 [=====] - 3s 57us/step - loss: 0.0132 - acc: 0.9959
- val_loss: 0.0861 - val_acc: 0.9777

Epoch 12/20
60000/60000 [=====] - 3s 57us/step - loss: 0.0132 - acc: 0.9957

```

- val_loss: 0.0826 - val_acc: 0.9777
Epoch 13/20
60000/60000 [=====] - 4s 61us/step - loss: 0.0096 - acc: 0.9971
- val_loss: 0.1023 - val_acc: 0.9731
Epoch 14/20
60000/60000 [=====] - 4s 61us/step - loss: 0.0103 - acc: 0.9966
- val_loss: 0.0762 - val_acc: 0.9803
Epoch 15/20
60000/60000 [=====] - 3s 56us/step - loss: 0.0079 - acc: 0.9976
- val_loss: 0.0845 - val_acc: 0.9804
Epoch 16/20
60000/60000 [=====] - 4s 60us/step - loss: 0.0074 - acc: 0.9975
- val_loss: 0.1045 - val_acc: 0.9756
Epoch 17/20
60000/60000 [=====] - 4s 62us/step - loss: 0.0092 - acc: 0.9970
- val_loss: 0.0873 - val_acc: 0.9767
Epoch 18/20
60000/60000 [=====] - 3s 58us/step - loss: 0.0084 - acc: 0.9972
- val_loss: 0.0867 - val_acc: 0.9782
Epoch 19/20
60000/60000 [=====] - 4s 63us/step - loss: 0.0089 - acc: 0.9970
- val_loss: 0.0911 - val_acc: 0.9778
Epoch 20/20
60000/60000 [=====] - 4s 61us/step - loss: 0.0060 - acc: 0.9981
- val_loss: 0.0757 - val_acc: 0.9806

```

In [59]:

```

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

score = model_batch.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ;
ax.set_ylabel('Categorical Crossentropy Loss');


# list of epoch numbers
x = list(range(1, nb_epoch+1))


vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.07565803507223973

Test accuracy: 0.9806





In [84]:


```
w_after = model_drop.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
plt.close()
```



In [0]:

In [0]:

Batch Normalization + Dropout

In [61]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras

from keras.layers import Dropout

model3 = Sequential()

model3.add(Dense(352, activation='relu', input_shape=(input_dim,), kernel_initializer=Ran
```

```

domNormal(mean=0.0, stddev=0.039, seed=None)))
model3.add(BatchNormalization())
model3.add(Dropout(0.5))

model3.add(Dense(52, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model3.add(BatchNormalization())
model3.add(Dropout(0.5))

model3.add(Dense(output_dim, activation='softmax'))

model3.summary()

```

Model: "sequential_22"

Layer (type)	Output Shape	Param #
dense_62 (Dense)	(None, 352)	276320
batch_normalization_7 (Batch Normalization)	(None, 352)	1408
dropout_5 (Dropout)	(None, 352)	0
dense_63 (Dense)	(None, 52)	18356
batch_normalization_8 (Batch Normalization)	(None, 52)	208
dropout_6 (Dropout)	(None, 52)	0
dense_64 (Dense)	(None, 10)	530
Total params: 296,822		
Trainable params: 296,014		
Non-trainable params: 808		

In [62]:

```

model3.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model3.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1,
, validation_data=(X_test, Y_test))

```

Train on 60000 samples, validate on 10000 samples

```

Epoch 1/20
60000/60000 [=====] - 5s 89us/step - loss: 0.5979 - acc: 0.8192
- val_loss: 0.1928 - val_acc: 0.9417
Epoch 2/20
60000/60000 [=====] - 4s 61us/step - loss: 0.3125 - acc: 0.9085
- val_loss: 0.1482 - val_acc: 0.9546
Epoch 3/20
60000/60000 [=====] - 4s 66us/step - loss: 0.2578 - acc: 0.9251
- val_loss: 0.1200 - val_acc: 0.9635
Epoch 4/20
60000/60000 [=====] - 4s 59us/step - loss: 0.2192 - acc: 0.9368
- val_loss: 0.1104 - val_acc: 0.9688
Epoch 5/20
60000/60000 [=====] - 3s 57us/step - loss: 0.1983 - acc: 0.9427
- val_loss: 0.0974 - val_acc: 0.9715
Epoch 6/20
60000/60000 [=====] - 4s 63us/step - loss: 0.1774 - acc: 0.9498
- val_loss: 0.0945 - val_acc: 0.9706
Epoch 7/20
60000/60000 [=====] - 3s 56us/step - loss: 0.1635 - acc: 0.9526
- val_loss: 0.0923 - val_acc: 0.9713
Epoch 8/20
60000/60000 [=====] - 4s 59us/step - loss: 0.1518 - acc: 0.9556
- val_loss: 0.0861 - val_acc: 0.9737
Epoch 9/20

```

```

60000/60000 [=====] - 4s 65us/step - loss: 0.1465 - acc: 0.9576
- val_loss: 0.0813 - val_acc: 0.9757
Epoch 10/20
60000/60000 [=====] - 4s 61us/step - loss: 0.1367 - acc: 0.9603
- val_loss: 0.0840 - val_acc: 0.9755
Epoch 11/20
60000/60000 [=====] - 4s 60us/step - loss: 0.1290 - acc: 0.9627
- val_loss: 0.0866 - val_acc: 0.9748
Epoch 12/20
60000/60000 [=====] - 4s 61us/step - loss: 0.1237 - acc: 0.9642
- val_loss: 0.0786 - val_acc: 0.9757
Epoch 13/20
60000/60000 [=====] - 4s 60us/step - loss: 0.1177 - acc: 0.9649
- val_loss: 0.0797 - val_acc: 0.9781
Epoch 14/20
60000/60000 [=====] - 4s 61us/step - loss: 0.1120 - acc: 0.9663
- val_loss: 0.0774 - val_acc: 0.9775
Epoch 15/20
60000/60000 [=====] - 4s 58us/step - loss: 0.1108 - acc: 0.9671
- val_loss: 0.0731 - val_acc: 0.9782
Epoch 16/20
60000/60000 [=====] - 4s 60us/step - loss: 0.1037 - acc: 0.9692
- val_loss: 0.0702 - val_acc: 0.9794
Epoch 17/20
60000/60000 [=====] - 4s 60us/step - loss: 0.0999 - acc: 0.9703
- val_loss: 0.0710 - val_acc: 0.9794
Epoch 18/20
60000/60000 [=====] - 4s 63us/step - loss: 0.0981 - acc: 0.9704
- val_loss: 0.0694 - val_acc: 0.9789
Epoch 19/20
60000/60000 [=====] - 3s 58us/step - loss: 0.0940 - acc: 0.9723
- val_loss: 0.0730 - val_acc: 0.9797
Epoch 20/20
60000/60000 [=====] - 4s 63us/step - loss: 0.0900 - acc: 0.9736
- val_loss: 0.0732 - val_acc: 0.9790

```

In [63]:

```

score = model3.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.07315509025455104
Test accuracy: 0.979





In [64]:


```
w_after = model3.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



In [0]:

In [0]:

Three Hidden Layer Architecture

Using RELU Activation and Adam Optimizer

In [0]:

In [65]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras

from keras.layers import Dropout

model1 = Sequential()

model1.add(Dense(352, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))

model1.add(Dense(52, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )

model1.add(Dense(102, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.6, seed=None)) )

model1.add(Dense(output_dim, activation='softmax'))

model1.summary()
```

Model: "sequential_23"

Layer (type)	Output Shape	Param #
dense_65 (Dense)	(None, 352)	276320
dense_66 (Dense)	(None, 52)	18356
dense_67 (Dense)	(None, 102)	5406
dense_68 (Dense)	(None, 10)	1030
Total params: 301,112		
Trainable params: 301,112		
Non-trainable params: 0		

In [66]:

```
model1.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model1.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 4s 63us/step - loss: 0.3633 - acc: 0.9046
- val_loss: 0.1766 - val_acc: 0.9475

Epoch 2/20

60000/60000 [=====] - 2s 39us/step - loss: 0.1396 - acc: 0.9590
- val_loss: 0.1287 - val_acc: 0.9616

Epoch 3/20

60000/60000 [=====] - 3s 42us/step - loss: 0.0965 - acc: 0.9714
- val_loss: 0.1168 - val_acc: 0.9675

Epoch 4/20

60000/60000 [=====] - 2s 41us/step - loss: 0.0773 - acc: 0.9760
- val_loss: 0.1324 - val_acc: 0.9637

Epoch 5/20

60000/60000 [=====] - 3s 42us/step - loss: 0.0646 - acc: 0.9796
- val_loss: 0.1348 - val_acc: 0.9646

Epoch 6/20

60000/60000 [=====] - 3s 43us/step - loss: 0.0517 - acc: 0.9835


```

- val_loss: 0.1337 - val_acc: 0.9676
Epoch 7/20
60000/60000 [=====] - 3s 44us/step - loss: 0.0440 - acc: 0.9863
- val_loss: 0.1049 - val_acc: 0.9724
Epoch 8/20
60000/60000 [=====] - 2s 39us/step - loss: 0.0380 - acc: 0.9880
- val_loss: 0.1191 - val_acc: 0.9725
Epoch 9/20
60000/60000 [=====] - 2s 39us/step - loss: 0.0377 - acc: 0.9880
- val_loss: 0.1243 - val_acc: 0.9726
Epoch 10/20
60000/60000 [=====] - 2s 40us/step - loss: 0.0356 - acc: 0.9892
- val_loss: 0.1401 - val_acc: 0.9703
Epoch 11/20
60000/60000 [=====] - 3s 42us/step - loss: 0.0305 - acc: 0.9908
- val_loss: 0.1245 - val_acc: 0.9731
Epoch 12/20
60000/60000 [=====] - 2s 38us/step - loss: 0.0302 - acc: 0.9907
- val_loss: 0.1260 - val_acc: 0.9717
Epoch 13/20
60000/60000 [=====] - 3s 42us/step - loss: 0.0236 - acc: 0.9926
- val_loss: 0.1347 - val_acc: 0.9727
Epoch 14/20
60000/60000 [=====] - 2s 40us/step - loss: 0.0180 - acc: 0.9943
- val_loss: 0.1305 - val_acc: 0.9731
Epoch 15/20
60000/60000 [=====] - 2s 41us/step - loss: 0.0325 - acc: 0.9904
- val_loss: 0.1257 - val_acc: 0.9747
Epoch 16/20
60000/60000 [=====] - 2s 40us/step - loss: 0.0245 - acc: 0.9924
- val_loss: 0.1310 - val_acc: 0.9757
Epoch 17/20
60000/60000 [=====] - 2s 38us/step - loss: 0.0167 - acc: 0.9946
- val_loss: 0.1460 - val_acc: 0.9708
Epoch 18/20
60000/60000 [=====] - 2s 39us/step - loss: 0.0168 - acc: 0.9949
- val_loss: 0.1398 - val_acc: 0.9743
Epoch 19/20
60000/60000 [=====] - 2s 41us/step - loss: 0.0209 - acc: 0.9938
- val_loss: 0.1822 - val_acc: 0.9692
Epoch 20/20
60000/60000 [=====] - 2s 41us/step - loss: 0.0236 - acc: 0.9928
- val_loss: 0.1301 - val_acc: 0.9767

```

In [67]:

```

score = model1.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']

```

```
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.13014392414744216
Test accuracy: 0.9767



In [68]:


```
w_after = model1.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



In [0]:

In [0]:

In [0]:

With Dropout

In [69]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras

from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(352, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_drop.add(Dropout(0.5))

model_drop.add(Dense(52, activation='relu', kernel_initializer=RandomNormal(mean=0.0, std dev=0.55, seed=None)) )
model_drop.add(Dropout(0.5))

model1.add(Dense(102, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.6, seed=None)) )
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()
```

Model: "sequential_24"

Layer (type)	Output Shape	Param #
=====		
dense_69 (Dense)	(None, 352)	276320
<hr/>		
dropout_7 (Dropout)	(None, 352)	0
<hr/>		
dense_70 (Dense)	(None, 52)	18356
<hr/>		
dropout_8 (Dropout)	(None, 52)	0
<hr/>		
dropout_9 (Dropout)	(None, 52)	0
<hr/>		
dense_72 (Dense)	(None, 10)	530
=====		
Total params: 295,206		
Trainable params: 295,206		
Non-trainable params: 0		

In [70]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 4s 74us/step - loss: 1.5708 - acc: 0.4904
- val_loss: 0.3985 - val_acc: 0.9120
Epoch 2/20
60000/60000 [=====] - 4s 74us/step - loss: 0.3985 - acc: 0.9120

```

60000/60000 [=====] - 3s 42us/step - loss: 0.8458 - acc: 0.7161
- val_loss: 0.2862 - val_acc: 0.9282
Epoch 3/20
60000/60000 [=====] - 2s 41us/step - loss: 0.6915 - acc: 0.7740
- val_loss: 0.2436 - val_acc: 0.9348
Epoch 4/20
60000/60000 [=====] - 2s 41us/step - loss: 0.5985 - acc: 0.8076
- val_loss: 0.2135 - val_acc: 0.9435
Epoch 5/20
60000/60000 [=====] - 3s 43us/step - loss: 0.5376 - acc: 0.8304
- val_loss: 0.2173 - val_acc: 0.9464
Epoch 6/20
60000/60000 [=====] - 3s 42us/step - loss: 0.5022 - acc: 0.8433
- val_loss: 0.1925 - val_acc: 0.9527
Epoch 7/20
60000/60000 [=====] - 2s 40us/step - loss: 0.4663 - acc: 0.8554
- val_loss: 0.1797 - val_acc: 0.9539
Epoch 8/20
60000/60000 [=====] - 3s 44us/step - loss: 0.4336 - acc: 0.8655
- val_loss: 0.1726 - val_acc: 0.9543
Epoch 9/20
60000/60000 [=====] - 3s 44us/step - loss: 0.4093 - acc: 0.8730
- val_loss: 0.1612 - val_acc: 0.9604
Epoch 10/20
60000/60000 [=====] - 3s 43us/step - loss: 0.3952 - acc: 0.8772
- val_loss: 0.1597 - val_acc: 0.9603
Epoch 11/20
60000/60000 [=====] - 2s 40us/step - loss: 0.3762 - acc: 0.8839
- val_loss: 0.1507 - val_acc: 0.9640
Epoch 12/20
60000/60000 [=====] - 2s 41us/step - loss: 0.3696 - acc: 0.8870
- val_loss: 0.1475 - val_acc: 0.9633
Epoch 13/20
60000/60000 [=====] - 2s 41us/step - loss: 0.3523 - acc: 0.8898
- val_loss: 0.1421 - val_acc: 0.9658
Epoch 14/20
60000/60000 [=====] - 2s 40us/step - loss: 0.3429 - acc: 0.8965
- val_loss: 0.1406 - val_acc: 0.9657
Epoch 15/20
60000/60000 [=====] - 2s 39us/step - loss: 0.3285 - acc: 0.8975
- val_loss: 0.1440 - val_acc: 0.9649
Epoch 16/20
60000/60000 [=====] - 3s 48us/step - loss: 0.3209 - acc: 0.9019
- val_loss: 0.1418 - val_acc: 0.9668
Epoch 17/20
60000/60000 [=====] - 3s 42us/step - loss: 0.3132 - acc: 0.9041
- val_loss: 0.1299 - val_acc: 0.9703
Epoch 18/20
60000/60000 [=====] - 2s 39us/step - loss: 0.3103 - acc: 0.9060
- val_loss: 0.1339 - val_acc: 0.9683
Epoch 19/20
60000/60000 [=====] - 2s 41us/step - loss: 0.3021 - acc: 0.9073
- val_loss: 0.1277 - val_acc: 0.9683
Epoch 20/20
60000/60000 [=====] - 2s 40us/step - loss: 0.2925 - acc: 0.9094
- val_loss: 0.1372 - val_acc: 0.9668

```

In [71]:

```

score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verb

```

```
ose=1, validation_data=(X_test, Y_test))
```


```
# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy
```

```
# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs
```

```
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.1372257164807059

Test accuracy: 0.9668



In [72]:


```
w_after = model_drop.get_weights()
```


```
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)
```

```
fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')
```

```
plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')
```

```
plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```





In [0]:

With Batch Normalization

In [0]:

In [73]:

```
# Multilayer perceptron

# https://intoli.com/blog/neural-network-initialization/
# If we sample weights from a normal distribution  $N(0,\sigma)$  we satisfy this condition with  $\sigma = \sqrt{2/(n_i+n_{i+1})}$ .
# h1 =>  $\sigma = \sqrt{2/(n_i+n_{i+1})} = 0.039 \Rightarrow N(0,\sigma) = N(0,0.039)$ 
# h2 =>  $\sigma = \sqrt{2/(n_i+n_{i+1})} = 0.055 \Rightarrow N(0,\sigma) = N(0,0.055)$ 
# h3 =>  $\sigma = \sqrt{2/(n_i+n_{i+1})} = 0.120 \Rightarrow N(0,\sigma) = N(0,0.120)$ 

from keras.layers.normalization import BatchNormalization

model_batch = Sequential()

model_batch.add(Dense(352, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(52, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(102, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))

model_batch.summary()
```

Model: "sequential_25"

Layer (type)	Output Shape	Param #
dense_73 (Dense)	(None, 352)	276320
batch_normalization_9 (Batch Normalization)	(None, 352)	1408
dense_74 (Dense)	(None, 52)	18356
batch_normalization_10 (Batch Normalization)	(None, 52)	208
dense_75 (Dense)	(None, 102)	5406
batch_normalization_11 (Batch Normalization)	(None, 102)	408
dense_76 (Dense)	(None, 10)	1030
Total params: 303,136		
Trainable params: 302,124		
Non-trainable params: 1,012		

In [74]:

```
model_batch.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model_batch.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 7s 109us/step - loss: 0.2560 - acc: 0.9258

- val_loss: 0.1237 - val_acc: 0.9627

Epoch 2/20

60000/60000 [=====] - 4s 72us/step - loss: 0.0924 - acc: 0.9721

- val_loss: 0.1167 - val_acc: 0.9627

Epoch 3/20

60000/60000 [=====] - 4s 71us/step - loss: 0.0614 - acc: 0.9813

- val_loss: 0.1097 - val_acc: 0.9666

Epoch 4/20

60000/60000 [=====] - 4s 70us/step - loss: 0.0465 - acc: 0.9855

- val_loss: 0.0835 - val_acc: 0.9754

Epoch 5/20

60000/60000 [=====] - 4s 66us/step - loss: 0.0369 - acc: 0.9876

- val_loss: 0.0784 - val_acc: 0.9750

Epoch 6/20

60000/60000 [=====] - 4s 73us/step - loss: 0.0313 - acc: 0.9897

- val_loss: 0.0795 - val_acc: 0.9753

Epoch 7/20

60000/60000 [=====] - 5s 78us/step - loss: 0.0253 - acc: 0.9918

- val_loss: 0.0890 - val_acc: 0.9735

Epoch 8/20

60000/60000 [=====] - 4s 73us/step - loss: 0.0231 - acc: 0.9925

- val_loss: 0.0789 - val_acc: 0.9787

Epoch 9/20

60000/60000 [=====] - 4s 70us/step - loss: 0.0183 - acc: 0.9942

- val_loss: 0.0930 - val_acc: 0.9735

Epoch 10/20

60000/60000 [=====] - 4s 70us/step - loss: 0.0174 - acc: 0.9940

- val_loss: 0.0863 - val_acc: 0.9759

Epoch 11/20

60000/60000 [=====] - 4s 74us/step - loss: 0.0159 - acc: 0.9945

- val_loss: 0.0812 - val_acc: 0.9790

Epoch 12/20

60000/60000 [=====] - 4s 69us/step - loss: 0.0141 - acc: 0.9953

- val_loss: 0.0953 - val_acc: 0.9759

Epoch 13/20

60000/60000 [=====] - 4s 71us/step - loss: 0.0147 - acc: 0.9951

- val_loss: 0.0896 - val_acc: 0.9769

Epoch 14/20

60000/60000 [=====] - 4s 71us/step - loss: 0.0130 - acc: 0.9957

- val_loss: 0.0832 - val_acc: 0.9799

Epoch 15/20

60000/60000 [=====] - 4s 68us/step - loss: 0.0128 - acc: 0.9961

- val_loss: 0.0938 - val_acc: 0.9764

Epoch 16/20

60000/60000 [=====] - 4s 72us/step - loss: 0.0093 - acc: 0.9973

- val_loss: 0.0813 - val_acc: 0.9786

Epoch 17/20

60000/60000 [=====] - 4s 73us/step - loss: 0.0111 - acc: 0.9964

- val_loss: 0.0976 - val_acc: 0.9755

Epoch 18/20

60000/60000 [=====] - 4s 70us/step - loss: 0.0103 - acc: 0.9964

- val_loss: 0.0868 - val_acc: 0.9791

Epoch 19/20

60000/60000 [=====] - 4s 69us/step - loss: 0.0100 - acc: 0.9965

- val_loss: 0.0964 - val_acc: 0.9784

Epoch 20/20

60000/60000 [=====] - 4s 69us/step - loss: 0.0088 - acc: 0.9971

- val_loss: 0.0903 - val_acc: 0.9787

In [75]:

```
# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

score = model_batch.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])


fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ;
ax.set_ylabel('Categorical Crossentropy Loss');

# list of epoch numbers
x = list(range(1, nb_epoch+1))

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.09027723879703226

Test accuracy: 0.9787



In [76]:

```
w_after = model_drop.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w, color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2')
```



```
plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
plt.close()
```



```
In [0]:
```

```
In [0]:
```

Batch Normalization + Dropout

```
In [77]:
```

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras

from keras.layers import Dropout

model3 = Sequential()

model3.add(Dense(352, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model3.add(BatchNormalization())
model3.add(Dropout(0.5))

model3.add(Dense(52, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model3.add(BatchNormalization())
model3.add(Dropout(0.5))

model3.add(Dense(102, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model3.add(BatchNormalization())
model3.add(Dropout(0.5))

model3.add(Dense(output_dim, activation='softmax'))

model3.summary()
```

Model: "sequential_26"

Layer (type)	Output Shape	Param #
dense_77 (Dense)	(None, 352)	276320
dense_78 (Dense)	(None, 52)	27128
dense_79 (Dense)	(None, 102)	103122
dense_80 (Dense)	(None, 10)	1020

batch_normalization_12 (Batch Normalization)	(None, 352)	1408
dropout_10 (Dropout)	(None, 352)	0
dense_78 (Dense)	(None, 52)	18356
batch_normalization_13 (Batch Normalization)	(None, 52)	208
dropout_11 (Dropout)	(None, 52)	0
dense_79 (Dense)	(None, 102)	5406
batch_normalization_14 (Batch Normalization)	(None, 102)	408
dropout_12 (Dropout)	(None, 102)	0
dense_80 (Dense)	(None, 10)	1030
=====		
Total params: 303,136		
Trainable params: 302,124		
Non-trainable params: 1,012		

In [78]:

```
model3.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model3.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1,
                    validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 7s 116us/step - loss: 0.9735 - acc: 0.6879
- val_loss: 0.2488 - val_acc: 0.9267
Epoch 2/20
60000/60000 [=====] - 4s 75us/step - loss: 0.4623 - acc: 0.8618
- val_loss: 0.1765 - val_acc: 0.9454
Epoch 3/20
60000/60000 [=====] - 5s 78us/step - loss: 0.3619 - acc: 0.8964
- val_loss: 0.1529 - val_acc: 0.9550
Epoch 4/20
60000/60000 [=====] - 5s 78us/step - loss: 0.3122 - acc: 0.9128
- val_loss: 0.1309 - val_acc: 0.9603
Epoch 5/20
60000/60000 [=====] - 4s 73us/step - loss: 0.2734 - acc: 0.9225
- val_loss: 0.1171 - val_acc: 0.9671
Epoch 6/20
60000/60000 [=====] - 4s 73us/step - loss: 0.2464 - acc: 0.9304
- val_loss: 0.1163 - val_acc: 0.9658
Epoch 7/20
60000/60000 [=====] - 4s 70us/step - loss: 0.2261 - acc: 0.9372
- val_loss: 0.1132 - val_acc: 0.9670
Epoch 8/20
60000/60000 [=====] - 5s 77us/step - loss: 0.2128 - acc: 0.9397
- val_loss: 0.1087 - val_acc: 0.9709
Epoch 9/20
60000/60000 [=====] - 4s 71us/step - loss: 0.1991 - acc: 0.9437
- val_loss: 0.0964 - val_acc: 0.9732
Epoch 10/20
60000/60000 [=====] - 4s 74us/step - loss: 0.1903 - acc: 0.9469
- val_loss: 0.0937 - val_acc: 0.9731
Epoch 11/20
60000/60000 [=====] - 4s 71us/step - loss: 0.1773 - acc: 0.9502
- val_loss: 0.0951 - val_acc: 0.9741
Epoch 12/20
60000/60000 [=====] - 4s 74us/step - loss: 0.1682 - acc: 0.9536
- val_loss: 0.0923 - val_acc: 0.9747
Epoch 13/20
60000/60000 [=====] - 5s 78us/step - loss: 0.1556 - acc: 0.9568
- val_loss: 0.0845 - val_acc: 0.9756
Epoch 14/20
60000/60000 [=====] - 4s 71us/step - loss: 0.1551 - acc: 0.9568
```

```

- val_loss: 0.0906 - val_acc: 0.9748
Epoch 15/20
60000/60000 [=====] - 5s 76us/step - loss: 0.1458 - acc: 0.9595
- val_loss: 0.0898 - val_acc: 0.9771
Epoch 16/20
60000/60000 [=====] - 4s 75us/step - loss: 0.1416 - acc: 0.9611
- val_loss: 0.0830 - val_acc: 0.9779
Epoch 17/20
60000/60000 [=====] - 5s 77us/step - loss: 0.1376 - acc: 0.9626
- val_loss: 0.0852 - val_acc: 0.9775
Epoch 18/20
60000/60000 [=====] - 4s 71us/step - loss: 0.1312 - acc: 0.9634
- val_loss: 0.0850 - val_acc: 0.9772
Epoch 19/20
60000/60000 [=====] - 5s 79us/step - loss: 0.1302 - acc: 0.9644
- val_loss: 0.0773 - val_acc: 0.9788
Epoch 20/20
60000/60000 [=====] - 5s 78us/step - loss: 0.1235 - acc: 0.9659
- val_loss: 0.0802 - val_acc: 0.9792

```

In [79]:

```

score = model3.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))


# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.08018281710293376
Test accuracy: 0.9792



In [80]:


```
w after = model3.get_weights()
```

```
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)
```

```
fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



In [0]:

In [0]:

In [0]:

Five Layer Architecture

Using RELU Activation and Adam Optimizer

In [0]:

In [91]:

```
model_relu = Sequential()
model_relu.add(Dense(250, activation='relu', input_shape=(input_dim,), kernel_initializer
=RandomNormal(mean=0.0, stddev=0.062, seed=None)))
```

```
model_relu.add(Dense(150, activation='relu', kernel_initializer=RandomNormal(mean=0.0, st
ddev=0.125, seed=None)) )
model_relu.add(Dense(146, activation='relu', kernel_initializer=RandomNormal(mean=0.0, st
ddev=0.15, seed=None)) )
model_relu.add(Dense(60, activation='relu', kernel_initializer=RandomNormal(mean=0.0, std
dev=0.25, seed=None)) )
model_relu.add(Dense(40, activation='relu', kernel_initializer=RandomNormal(mean=0.0, std
dev=0.5, seed=None)) )
model_relu.add(Dense(output_dim, activation='softmax'))

print(model_relu.summary())
```

Model: "sequential_29"

Layer (type)	Output Shape	Param #
dense_93 (Dense)	(None, 250)	196250
dense_94 (Dense)	(None, 150)	37650
dense_95 (Dense)	(None, 146)	22046
dense_96 (Dense)	(None, 60)	8820
dense_97 (Dense)	(None, 40)	2440
dense_98 (Dense)	(None, 10)	410
Total params: 267,616		
Trainable params: 267,616		
Non-trainable params: 0		
None		

In [0]:

In [92]:

```
model_relu.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy
'])

history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbo
se=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20
60000/60000 [=====] - 5s 90us/step - loss: 0.3360 - acc: 0.9028
- val_loss: 0.1423 - val_acc: 0.9570

Epoch 2/20
60000/60000 [=====] - 3s 47us/step - loss: 0.1223 - acc: 0.9639
- val_loss: 0.1142 - val_acc: 0.9638

Epoch 3/20
60000/60000 [=====] - 3s 45us/step - loss: 0.0848 - acc: 0.9738
- val_loss: 0.0979 - val_acc: 0.9721

Epoch 4/20
60000/60000 [=====] - 3s 45us/step - loss: 0.0639 - acc: 0.9802
- val_loss: 0.1029 - val_acc: 0.9688

Epoch 5/20
60000/60000 [=====] - 3s 46us/step - loss: 0.0529 - acc: 0.9832
- val_loss: 0.0985 - val_acc: 0.9701

Epoch 6/20
60000/60000 [=====] - 3s 46us/step - loss: 0.0453 - acc: 0.9850
- val_loss: 0.0847 - val_acc: 0.9762

Epoch 7/20
60000/60000 [=====] - 3s 47us/step - loss: 0.0388 - acc: 0.9875
- val_loss: 0.0940 - val_acc: 0.9746

Epoch 8/20
60000/60000 [=====] - 3s 47us/step - loss: 0.0314 - acc: 0.9896
- val_loss: 0.1085 - val_acc: 0.9718

Epoch 9/20

```

60000/60000 [=====] - 3s 52us/step - loss: 0.0338 - acc: 0.9891
- val_loss: 0.0975 - val_acc: 0.9740
Epoch 10/20
60000/60000 [=====] - 3s 47us/step - loss: 0.0282 - acc: 0.9908
- val_loss: 0.1011 - val_acc: 0.9754
Epoch 11/20
60000/60000 [=====] - 3s 44us/step - loss: 0.0249 - acc: 0.9920
- val_loss: 0.1265 - val_acc: 0.9700
Epoch 12/20
60000/60000 [=====] - 3s 49us/step - loss: 0.0253 - acc: 0.9918
- val_loss: 0.0944 - val_acc: 0.9781
Epoch 13/20
60000/60000 [=====] - 3s 48us/step - loss: 0.0233 - acc: 0.9924
- val_loss: 0.0943 - val_acc: 0.9780
Epoch 14/20
60000/60000 [=====] - 3s 47us/step - loss: 0.0223 - acc: 0.9929
- val_loss: 0.0967 - val_acc: 0.9768
Epoch 15/20
60000/60000 [=====] - 3s 49us/step - loss: 0.0175 - acc: 0.9942
- val_loss: 0.1023 - val_acc: 0.9771
Epoch 16/20
60000/60000 [=====] - 3s 48us/step - loss: 0.0166 - acc: 0.9946
- val_loss: 0.0940 - val_acc: 0.9784
Epoch 17/20
60000/60000 [=====] - 3s 46us/step - loss: 0.0159 - acc: 0.9951
- val_loss: 0.1080 - val_acc: 0.9759
Epoch 18/20
60000/60000 [=====] - 3s 49us/step - loss: 0.0198 - acc: 0.9937
- val_loss: 0.1146 - val_acc: 0.9758
Epoch 19/20
60000/60000 [=====] - 3s 50us/step - loss: 0.0127 - acc: 0.9959
- val_loss: 0.1031 - val_acc: 0.9774
Epoch 20/20
60000/60000 [=====] - 3s 47us/step - loss: 0.0180 - acc: 0.9943
- val_loss: 0.1213 - val_acc: 0.9743

```

In [93]:

```

score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1, nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs


vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.12134343287702913

Test accuracy: 0.9743





In [94]:


```
w_after = model_relu.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
plt.close()
```



With Dropout

In [85]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras

from keras.layers import Dropout

model_drop = Sequential()
```

```
model_drop.add(Dense(250, activation='relu', input_shape=(input_dim,),kernel_initializer=
RandomNormal(mean=0.0, stddev=0.125, seed=None)))
model_drop.add(Dropout(0.5))

model_drop.add(Dense(150, activation='relu',
                      kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)
))
model_drop.add(Dropout(0.5))

model_drop.add(Dense(146, activation='relu',
                      kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)
))
model_drop.add(Dropout(0.5))
model_drop.add(Dense(60, activation='relu',
                      kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)
))
model_drop.add(Dropout(0.5))

model_drop.add(Dense(40, activation='relu',
                      kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, seed=None)
))
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))
print(model_drop.summary())
```

Model: "sequential_28"

Layer (type)	Output Shape	Param #
=====		
dense_87 (Dense)	(None, 250)	196250
dropout_13 (Dropout)	(None, 250)	0
dense_88 (Dense)	(None, 150)	37650
dropout_14 (Dropout)	(None, 150)	0
dense_89 (Dense)	(None, 146)	22046
dropout_15 (Dropout)	(None, 146)	0
dense_90 (Dense)	(None, 60)	8820
dropout_16 (Dropout)	(None, 60)	0
dense_91 (Dense)	(None, 40)	2440
dropout_17 (Dropout)	(None, 40)	0
dense_92 (Dense)	(None, 10)	410
=====		

Total params: 267,616
Trainable params: 267,616
Non-trainable params: 0

None

In [86]:

```
model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy
'])

history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbo
se=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 6s 92us/step - loss: 2.3843 - acc: 0.1383
- val loss: 1.8483 - val acc: 0.4532


```

Epoch 2/20
60000/60000 [=====] - 3s 58us/step - loss: 1.5475 - acc: 0.4169
- val_loss: 1.0906 - val_acc: 0.6451
Epoch 3/20
60000/60000 [=====] - 3s 54us/step - loss: 1.1712 - acc: 0.5643
- val_loss: 0.8176 - val_acc: 0.7626
Epoch 4/20
60000/60000 [=====] - 3s 52us/step - loss: 0.9148 - acc: 0.6899
- val_loss: 0.5543 - val_acc: 0.8483
Epoch 5/20
60000/60000 [=====] - 3s 51us/step - loss: 0.7204 - acc: 0.7777
- val_loss: 0.4719 - val_acc: 0.8741
Epoch 6/20
60000/60000 [=====] - 3s 49us/step - loss: 0.6096 - acc: 0.8303
- val_loss: 0.3595 - val_acc: 0.9196
Epoch 7/20
60000/60000 [=====] - 3s 51us/step - loss: 0.5343 - acc: 0.8612
- val_loss: 0.2917 - val_acc: 0.9379
Epoch 8/20
60000/60000 [=====] - 3s 52us/step - loss: 0.4666 - acc: 0.8813
- val_loss: 0.2648 - val_acc: 0.9429
Epoch 9/20
60000/60000 [=====] - 3s 58us/step - loss: 0.4232 - acc: 0.8957
- val_loss: 0.2428 - val_acc: 0.9481
Epoch 10/20
60000/60000 [=====] - 3s 53us/step - loss: 0.3869 - acc: 0.9049
- val_loss: 0.2383 - val_acc: 0.9479
Epoch 11/20
60000/60000 [=====] - 3s 50us/step - loss: 0.3681 - acc: 0.9101
- val_loss: 0.2254 - val_acc: 0.9515
Epoch 12/20
60000/60000 [=====] - 3s 52us/step - loss: 0.3393 - acc: 0.9177
- val_loss: 0.2145 - val_acc: 0.9558
Epoch 13/20
60000/60000 [=====] - 3s 50us/step - loss: 0.3302 - acc: 0.9222
- val_loss: 0.2015 - val_acc: 0.9576
Epoch 14/20
60000/60000 [=====] - 3s 54us/step - loss: 0.3135 - acc: 0.9253
- val_loss: 0.1920 - val_acc: 0.9582
Epoch 15/20
60000/60000 [=====] - 3s 55us/step - loss: 0.2979 - acc: 0.9295
- val_loss: 0.1886 - val_acc: 0.9617
Epoch 16/20
60000/60000 [=====] - 3s 52us/step - loss: 0.2828 - acc: 0.9327
- val_loss: 0.1806 - val_acc: 0.9624
Epoch 17/20
60000/60000 [=====] - 3s 52us/step - loss: 0.2742 - acc: 0.9351
- val_loss: 0.1774 - val_acc: 0.9638
Epoch 18/20
60000/60000 [=====] - 3s 51us/step - loss: 0.2672 - acc: 0.9374
- val_loss: 0.1788 - val_acc: 0.9645
Epoch 19/20
60000/60000 [=====] - 3s 51us/step - loss: 0.2624 - acc: 0.9383
- val_loss: 0.1715 - val_acc: 0.9653
Epoch 20/20
60000/60000 [=====] - 3s 53us/step - loss: 0.2526 - acc: 0.9402
- val_loss: 0.1637 - val_acc: 0.9662

```

In [88]:

```

score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())

```


```
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```

Test score: 0.16367745212987064
Test accuracy: 0.9662



In [90]:


```
w_after = model_drop.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```





In [0]:

In [0]:

In [0]:

With Batch Normalization

In [0]:

In [95]:

```
# Multilayer perceptron

# https://intoli.com/blog/neural-network-initialization/
# If we sample weights from a normal distribution  $N(0, \sigma)$  we satisfy this condition with  $\sigma = \sqrt{2/(n_i + n_{i+1})}$ .
# h1 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.039 \Rightarrow N(0, \sigma) = N(0, 0.039)$ 
# h2 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.055 \Rightarrow N(0, \sigma) = N(0, 0.055)$ 
# h1 =>  $\sigma = \sqrt{2/(n_i + n_{i+1})} = 0.120 \Rightarrow N(0, \sigma) = N(0, 0.120)$ 

from keras.layers.normalization import BatchNormalization

model_batch = Sequential()

model_batch.add(Dense(250, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(150, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(146, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.15, seed=None)))
model_batch.add(BatchNormalization())

model_drop.add(Dense(60, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.2, seed=None)))
model_batch.add(BatchNormalization())

model_drop.add(Dense(40, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.6, seed=None)))
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))

model_batch.summary()
```

Model: "sequential_30"

Layer (type)	Output Shape	Param #
--------------	--------------	---------

Layer (type)	Output shape	Param #
dense_99 (Dense)	(None, 250)	196250
batch_normalization_15 (Batch Normalization)	(None, 250)	1000
dense_100 (Dense)	(None, 150)	37650
batch_normalization_16 (Batch Normalization)	(None, 150)	600
dense_101 (Dense)	(None, 146)	22046
batch_normalization_17 (Batch Normalization)	(None, 146)	584
batch_normalization_18 (Batch Normalization)	(None, 146)	584
batch_normalization_19 (Batch Normalization)	(None, 146)	584
dense_104 (Dense)	(None, 10)	1470
Total params: 260,768		
Trainable params: 259,092		
Non-trainable params: 1,676		

In [96]:

```
model_batch.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_batch.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 9s 156us/step - loss: 0.2259 - acc: 0.9315
- val_loss: 0.1145 - val_acc: 0.9643
Epoch 2/20
60000/60000 [=====] - 6s 95us/step - loss: 0.0863 - acc: 0.9736
- val_loss: 0.0981 - val_acc: 0.9702
Epoch 3/20
60000/60000 [=====] - 6s 97us/step - loss: 0.0583 - acc: 0.9819
- val_loss: 0.0806 - val_acc: 0.9726
Epoch 4/20
60000/60000 [=====] - 6s 93us/step - loss: 0.0449 - acc: 0.9854
- val_loss: 0.1113 - val_acc: 0.9627
Epoch 5/20
60000/60000 [=====] - 6s 97us/step - loss: 0.0316 - acc: 0.9900
- val_loss: 0.0823 - val_acc: 0.9744
Epoch 6/20
60000/60000 [=====] - 6s 96us/step - loss: 0.0255 - acc: 0.9918
- val_loss: 0.0806 - val_acc: 0.9775
Epoch 7/20
60000/60000 [=====] - 6s 98us/step - loss: 0.0225 - acc: 0.9923
- val_loss: 0.0897 - val_acc: 0.9726
Epoch 8/20
60000/60000 [=====] - 6s 96us/step - loss: 0.0221 - acc: 0.9930
- val_loss: 0.0891 - val_acc: 0.9734
Epoch 9/20
60000/60000 [=====] - 6s 103us/step - loss: 0.0205 - acc: 0.9927
- val_loss: 0.0805 - val_acc: 0.9770
Epoch 10/20
60000/60000 [=====] - 6s 99us/step - loss: 0.0158 - acc: 0.9947
- val_loss: 0.0888 - val_acc: 0.9737
Epoch 11/20
60000/60000 [=====] - 6s 96us/step - loss: 0.0147 - acc: 0.9954
- val_loss: 0.0883 - val_acc: 0.9769
Epoch 12/20
60000/60000 [=====] - 6s 93us/step - loss: 0.0120 - acc: 0.9959
- val_loss: 0.0881 - val_acc: 0.9772
Epoch 13/20
60000/60000 [=====] - 6s 102us/step - loss: 0.0125 - acc: 0.9961
- val_loss: 0.0867 - val_acc: 0.9764
```

```

val_loss: 0.0862 - val_acc: 0.9792
Epoch 14/20
60000/60000 [=====] - 6s 95us/step - loss: 0.0166 - acc: 0.9942
- val_loss: 0.0892 - val_acc: 0.9759
Epoch 15/20
60000/60000 [=====] - 6s 94us/step - loss: 0.0122 - acc: 0.9958
- val_loss: 0.0832 - val_acc: 0.9788
Epoch 16/20
60000/60000 [=====] - 6s 94us/step - loss: 0.0098 - acc: 0.9968
- val_loss: 0.0823 - val_acc: 0.9788
Epoch 17/20
60000/60000 [=====] - 6s 97us/step - loss: 0.0088 - acc: 0.9971
- val_loss: 0.0844 - val_acc: 0.9784
Epoch 18/20
60000/60000 [=====] - 6s 93us/step - loss: 0.0073 - acc: 0.9975
- val_loss: 0.0985 - val_acc: 0.9755
Epoch 19/20
60000/60000 [=====] - 6s 96us/step - loss: 0.0115 - acc: 0.9963
- val_loss: 0.0845 - val_acc: 0.9784
Epoch 20/20
60000/60000 [=====] - 6s 101us/step - loss: 0.0090 - acc: 0.9970
- val_loss: 0.0862 - val_acc: 0.9792

```

In [97]:

```

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

score = model_batch.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])

fig, ax = plt.subplots(1, 1)
ax.set_xlabel('epoch') ;
ax.set_ylabel('Categorical Crossentropy Loss');


# list of epoch numbers
x = list(range(1, nb_epoch+1))

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.08622197609168743

Test accuracy: 0.9792



In [98]:


```
w_after = model_drop.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
plt.close()
```



In [0]:

In [0]:

Batch Normalization + Dropout

In [99]:

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in-keras

from keras.layers import Dropout

model3 = Sequential()

model3.add(Dense(250, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.039, seed=None)))
model3.add(BatchNormalization())
model3.add(Dropout(0.5))
```

```
model3.add(Dense(150, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.55, seed=None)) )
model3.add(BatchNormalization())
model3.add(Dropout(0.5))

model3.add(Dense(146, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.15, seed=None)) )
model3.add(BatchNormalization())
model3.add(Dropout(0.5))

model3.add(Dense(60, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.2, seed=None)) )
model3.add(BatchNormalization())
model3.add(Dropout(0.5))

model3.add(Dense(40, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.6, seed=None)) )
model3.add(BatchNormalization())
model3.add(Dropout(0.5))

model3.add(Dense(output_dim, activation='softmax'))

model3.summary()
```

Model: "sequential_31"

Layer (type)	Output Shape	Param #
=====		
dense_105 (Dense)	(None, 250)	196250
batch_normalization_20 (Batch Normalization)	(None, 250)	1000
dropout_18 (Dropout)	(None, 250)	0
dense_106 (Dense)	(None, 150)	37650
batch_normalization_21 (Batch Normalization)	(None, 150)	600
dropout_19 (Dropout)	(None, 150)	0
dense_107 (Dense)	(None, 146)	22046
batch_normalization_22 (Batch Normalization)	(None, 146)	584
dropout_20 (Dropout)	(None, 146)	0
dense_108 (Dense)	(None, 60)	8820
batch_normalization_23 (Batch Normalization)	(None, 60)	240
dropout_21 (Dropout)	(None, 60)	0
dense_109 (Dense)	(None, 40)	2440
batch_normalization_24 (Batch Normalization)	(None, 40)	160
dropout_22 (Dropout)	(None, 40)	0
dense_110 (Dense)	(None, 10)	410
=====		
Total params: 270,200		
Trainable params: 268,908		
Non-trainable params: 1,292		

In [100]:

```
model3.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model3.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1)
```

```
, validation_data=(X_test, Y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 11s 176us/step - loss: 1.8646 - acc: 0.3724 - val_loss: 0.6914 - val_acc: 0.8356

Epoch 2/20

60000/60000 [=====] - 7s 113us/step - loss: 0.9432 - acc: 0.6832 - val_loss: 0.3925 - val_acc: 0.9030

Epoch 3/20

60000/60000 [=====] - 6s 105us/step - loss: 0.6758 - acc: 0.7879 - val_loss: 0.2720 - val_acc: 0.9294

Epoch 4/20

60000/60000 [=====] - 7s 114us/step - loss: 0.5389 - acc: 0.8472 - val_loss: 0.2094 - val_acc: 0.9448

Epoch 5/20

60000/60000 [=====] - 7s 110us/step - loss: 0.4560 - acc: 0.8775 - val_loss: 0.1955 - val_acc: 0.9483

Epoch 6/20

60000/60000 [=====] - 6s 108us/step - loss: 0.4007 - acc: 0.8953 - val_loss: 0.1827 - val_acc: 0.9521

Epoch 7/20

60000/60000 [=====] - 7s 113us/step - loss: 0.3588 - acc: 0.9081 - val_loss: 0.1637 - val_acc: 0.9573

Epoch 8/20

60000/60000 [=====] - 7s 112us/step - loss: 0.3302 - acc: 0.9183 - val_loss: 0.1518 - val_acc: 0.9612

Epoch 9/20

60000/60000 [=====] - 7s 115us/step - loss: 0.3078 - acc: 0.9257 - val_loss: 0.1429 - val_acc: 0.9649

Epoch 10/20

60000/60000 [=====] - 6s 108us/step - loss: 0.2910 - acc: 0.9299 - val_loss: 0.1329 - val_acc: 0.9658

Epoch 11/20

60000/60000 [=====] - 7s 109us/step - loss: 0.2740 - acc: 0.9345 - val_loss: 0.1351 - val_acc: 0.9684

Epoch 12/20

60000/60000 [=====] - 7s 111us/step - loss: 0.2668 - acc: 0.9364 - val_loss: 0.1231 - val_acc: 0.9694

Epoch 13/20

60000/60000 [=====] - 7s 110us/step - loss: 0.2528 - acc: 0.9393 - val_loss: 0.1252 - val_acc: 0.9697

Epoch 14/20

60000/60000 [=====] - 7s 125us/step - loss: 0.2406 - acc: 0.9432 - val_loss: 0.1233 - val_acc: 0.9690

Epoch 15/20

60000/60000 [=====] - 7s 112us/step - loss: 0.2359 - acc: 0.9446 - val_loss: 0.1157 - val_acc: 0.9732

Epoch 16/20

60000/60000 [=====] - 7s 117us/step - loss: 0.2280 - acc: 0.9461 - val_loss: 0.1172 - val_acc: 0.9729

Epoch 17/20

60000/60000 [=====] - 6s 108us/step - loss: 0.2165 - acc: 0.9489 - val_loss: 0.1124 - val_acc: 0.9727

Epoch 18/20

60000/60000 [=====] - 7s 112us/step - loss: 0.2213 - acc: 0.9493 - val_loss: 0.1105 - val_acc: 0.9736

Epoch 19/20

60000/60000 [=====] - 7s 110us/step - loss: 0.2119 - acc: 0.9508 - val_loss: 0.1037 - val_acc: 0.9759

Epoch 20/20

60000/60000 [=====] - 7s 116us/step - loss: 0.1936 - acc: 0.9545 - val_loss: 0.1110 - val_acc: 0.9735

In [101]:

```
score = model3.evaluate(X_test, Y_test, verbose=0)
```

```
print('Test score:', score[0])
```

```
print('Test accuracy:', score[1])
```

```
fig,ax = plt.subplots(1,1)
```

```
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
```



```

# list of epoch numbers
x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy


# loss : training loss
# acc : train accuracy
# for each key in history.history we will have a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)

```

Test score: 0.11095908558527008

Test accuracy: 0.9735



In [102]:

```

w_after = model3.get_weights()


h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)


fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()

```





In [0]:

In [0]:

CONCLUSION

In [103]:

```
from prettytable import PrettyTable

x=PrettyTable()

x.field_names=(['No.of.Layers','Layers in MLP','Model Type','Test Score','Test accuracy'
])

x.add_row(['2-Layered','352-52','Without Dropout and BN',0.10,0.97])
x.add_row(['2-Layered','352-52','With Dropout',0.13,0.96])
x.add_row(['2-Layered','352-52','With BN',0.08,0.98])
x.add_row(['2-Layered','352-52','Dropout+BN',0.072,0.97])

print(x)
```

No.of.Layers	Layers in MLP	Model Type	Test Score	Test accuracy
2-Layered	352-52	Without Dropout and BN	0.1	0.97
2-Layered	352-52	With Dropout	0.13	0.96
2-Layered	352-52	With BN	0.08	0.98
2-Layered	352-52	Dropout+BN	0.072	0.97

In [104]:

```
y=PrettyTable()

y.field_names=(['No.of.Layers','Layers in MLP','Model Type','Test Score','Test Value'])

y.add_row(['3-Layered','352-52-102','Without Dropout and BN',0.15,0.975])
y.add_row(['3-Layered','352-52-102','With Dropout',0.14,0.967])
y.add_row(['3-Layered','352-52-102','With BN',0.09,0.979])
y.add_row(['3-Layered','352-52-102','Dropout+BN',0.07,0.978])

print(y)
```

No.of.Layers	Layers in MLP	Model Type	Test Score	Test Value
3-Layered	352-52-102	Without Dropout and BN	0.15	0.975
3-Layered	352-52-102	With Dropout	0.14	0.967
3-Layered	352-52-102	With BN	0.09	0.979
3-Layered	352-52-102	Dropout+BN	0.07	0.978

```
In [108]:
```

```
z=PrettyTable()

z.field_names=(['No.of.Layers','Layers in MLP','Model Type','Test Score','Test Value'])

z.add_row(['5-Layered','250-150-146-60-40','With Dropout and BN',0.12,0.974])
z.add_row(['5-Layered','250-150-146-60-40','With Dropout',0.16,0.966])
z.add_row(['5-Layered','250-150-146-60-40','With BN',0.08,0.979])
z.add_row(['5-Layered','250-150-146-60-40','Dropout+BN',0.11,0.973])

print(z)
```

No.of.Layers	Layers in MLP	Model Type	Test Score	Test Value
5-Layered	250-150-146-60-40	With Dropout and BN	0.12	0.974
5-Layered	250-150-146-60-40	With Dropout	0.16	0.966
5-Layered	250-150-146-60-40	With BN	0.08	0.979
5-Layered	250-150-146-60-40	Dropout+BN	0.11	0.973

1)Here we have use Mutli-Layered perceptrons Architecture where we have used 2-Layered, 3-Layered and 5-Layered Structures, with different number of neurons.

2)By using 2-layered Architecture with 352-52 neurons, we have seen that by adding Batch Normalization gave highest accuracy.

3)By using 3-Layered Architecture with 352-52-102 neurons, we have seen that batch normalization and both droput and batch normalization gave highest accuracy

4)By using 5-Layered Architecture with 250-150-146-60-40 neurons, adding batch normalization gave highest accuracy.

```
In [0]:
```

```
In [0]:
```